

Advanced Fortran Topics

Partly a



Event

Reinhold Bader Gilbert Brietzke

Leibniz Supercomputing Centre Munich

Fortran features under consideration



Continuing Standardization process:

	Fortran 66	ancient	
	Fortran 77 (1980)	traditional	
F95	Fortran 90 (1991)	large revision	
	Fortran 95 (1997)	small revision	
F03	Fortran 2003 (2004)	large revision	
F08	Fortran 2008 (2010)	mid-size revision	
	TS 29113 (2012)	extends C interop	integration into 🕫
	TS 18508 (2015)	extends parallelism	is completed
F18	Fortran 2018 (2018)	FDIS out for vote	-
			-

- Focus of this course is on Fos and Fos
- - the two Technical Specifications will also be (partially) covered

Overview of covered features



- Day 1:
 - recapitulation of important features; object-based programming. First steps into object orientation
- Day 2:
 - further object-oriented features, I/O extensions, IEEE FP processing
- Day 3:
 - generic features, interoperation with C
- Day 4 "Design Patterns"
 - how to use the OO features; first intro to PGAS programming
- Day 5 "PGAS":
 - parallel programming with coarrays
- Exercises: interspersed with talks see printed schedule
- Prerequisites:
 - good knowledge of F95
 - as covered e.g., in the winter event "Programming with Fortran" (and some own experience, if possible)
 - some knowledge of OpenMP (shared memory parallelism)

Social Event and Guided Tour



If desired by participants:

 joint dinner (self-funded) in the centre of Garching (Neuwirt) on Monday evening at 19:00

Guided Tour through the computer rooms at LRZ

- on Wednesday starting 18:00, approximately 60 minutes
- courtesy Volker Weinberg and Reinhold Bader



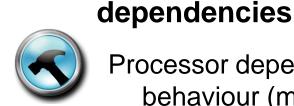
Conventions and Flags used in these talks



Standards conformance



Recommended practice



Processor dependent behaviour (may be unportable)

Implementation

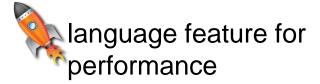


Standard conforming, but considered questionable style



Dangerous practice, likely to introduce bugs and/or non-conforming behaviour







Gotcha! Non-conforming and/or definitely buggy

Some references



- Modern Fortran explained (7th edition)
 - Michael Metcalf, John Reid, Malcolm Cohen, OUP, 2011
- The Fortran 2003 Handbook
 - J. Adams, W. Brainerd, R. Hendrickson, R. Maine, J. Martin, B. Smith.
 Springer, 2008
- Guide to Fortran 2008 programming (introductory text)
 - W. Brainerd. Springer, 2015
- Modern Fortran Style and Usage (best practices guide)
 - N. Clerman, W. Spector. Cambridge University Press, 2012
- Scientific Software Design The Object-Oriented Way (1st edition)
 - Damian Rouson, Jim Xia, Xiaofeng Xu, Cambridge, 2011

References cont'd



- Design Patterns Elements of Reusable Object-oriented Software
 - E. Gamma, R. Helm, R. Johnson, J. Vlissides. Addison-Wesley, 1994
- Modern Fortran in Practice
 - Arjen Markus, Cambridge University Press, 2012
- Introduction to High Performance Computing for Scientists and Engineers
 - G. Hager and G. Wellein
- Download of (updated) PDFs of the slides and exercise archive
 - freely available under a creative commons license
 - https://doku.lrz.de/display/PUBLIC/Materials+-+Advanced+Fortran+Topics



Recapitulation: Module features Global variables The environment problem

Some features from Fortran 2003

What is a module?



A program unit

... that permits packaging of

- procedure interfaces
- global variables
- named constants
- type definitions
- named interfaces
- procedure implementations

for reuse,

... as well as supporting

- information hiding
- (limited) namespace management

Module definition syntax

```
module <module-name>
    [ specification-part ]
contains
    [ module-subprogram, ...]
end module <module-name>
```

executable statements in a module can only appear in module subprograms

also known as encapsulation

Illustrative example: heat conduction in 2 dimensions



Simplest case:

 partial differential equation for temperature Φ(x, y, t)

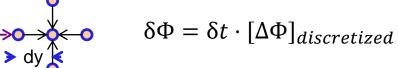
$$\frac{\partial \Phi}{\partial t} = \frac{\partial^2 \Phi}{\partial x^2} + \frac{\partial^2 \Phi}{\partial y^2}$$

- produce stationary solution on a (unit) square
- provided: initial values inside;
 boundary values on NSWE edges

w E → t=∞

Numeric model:

- discretize square: increments dx, dy
- real array phi(:,:) of rank 2 models temperature field
 → use a global variable for this
- iteration process (Jacobi)
 based on 5 point stencil
 generates stationary solution



with given time increment δt

Heat conduction: Global variable declarations



```
mod heat
phi
phinew
   ndim = 200
   dk = ...
   heat ival()
   heat_bval()
   heat_iter()
```

public entities: indicated

by yellow background

```
module mod heat
    implicit none
                            changes default accessibility
    private ---
    public :: heat ival, heat bval, heat iter
access to procedures
    integer, parameter, public :: ndim = 200
    integer, parameter, public :: dk = ...
                                            encapsulated
    real(dk) :: phi (ndim,ndim)
                                             global data
    real(dk) :: phinew (2:ndim-1,2:ndim-1)
  contains
    : ! implementation of module
     : ! procedures - see later slides
  end module mod heat
```

Call sequence:

- set initial and boundary values
- repeatedly iterate until convergence
- print out result

Module procedures for the heat fields (1)



Perform iteration steps

```
real(dk) function heat iter (dt, num)
  real(dk), intent(in) :: dt
  integer, intent(in) :: num
  real(dk) :: dphimax, dphi
  do it = 1, num
    dphimax = 0.0 dk
    do j = 2, size(phi, 2) - 1
      do i = 2, size(phi,1) - 1
         dphi = dt * (...) discretization needs phi
  on neighbouring points
         phinew(i, j) = phi(i, j) + dphi
         dphimax = max(dphi, dphimax)
      end do 1
                    dependency forces use
                    of auxiliary field phinew
   end do
   phi(2:size(phi,1)-1, &
        2:size(phi,2)-1) = phinew
  end do
  heat iter = dphimax
end function
```

preserves boundary values

Notes:

- heat_iter is a module function in mod_heat
 → it has access to fields phi, phinew by host association
- global variables declared in a module are persistent: they implicitly have the SAVE attribute

 An explicit SAVE can be

of a procedure.

specified for local variables

 example code provided as basis for future exercises

Module procedures for the heat fields (2): Procedure arguments



Example's boundary and initial value conditions:

- provided via functions
- a function or subroutine can be a dummy argument

F03 Abstract interface

in specification part of module

 describes interface of a procedure that does not (yet) exist

Initial value settings

module procedure heat_ival in mod heat

(Legacy) Alternative

 replace procedure statement by a (regular) interface block

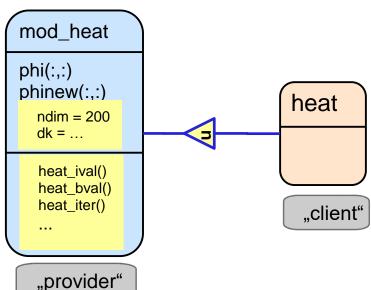
Heat main program



Implements call sequence

Graphical representation

```
program heat
  use mod heat
  use mystuff provides myfun etc.
  implicit none
  integer :: it
  real(dk), parameter :: &
            eps = 1e-6 dk
  real(dk) :: dt no use access
  phi(3,4) = 0.0 dk
  call heat ival(myfun)
  call heat bval(...)
                              OK, public
  do
                           module procedures
    dt = ... ! time step
    if (heat iter(dt, 1) < eps) exit
  end do
  :! print results
end program
```



use association provides an inheritance mechanism (for all public entities of a module)

Some



Fo3) extensions for handling of globals



Define entities which exist only once

example: temperature field

Analogous for derived types:

client should not be able to create entity of a type

```
module mod ptype
  implicit none
  private
  type, private :: ptype
    : ! type components
  end type
  type(ptype), public :: o_ptype
end module
                           public object
                          of private type
```

- type components can be public
- "Singleton" programming pattern

Client usage:

access public type components

```
type is private
use mod_ptype
                              → will not compile
type(ptype) :: o2
o ptype%i = 4
                              access to
                            singleton OK
```

"Read-only" objects:

```
type(ptype), public, &
          protected :: o_ptype
```

- attribute can not be applied to type definitions / components
- client outside defining module shall not define object (neither directly or indirectly e.g., via a pointer)
- for a pointer, PROTECTED refers to association status

Global entities: Threading issues



Typical threading model used

- OpenMP (assuming some knowledge here)
- directive based method for shared memory parallelism

Question discussed here:

- What happens if global variables need to be accessed from threaded parts of the code?
- How can "thread-safeness" be achieved?

Example: counting objects



```
module mod foo
  type :: foo
    private
    real, allocatable :: stuff(:)
  end type
  integer, protected, foo count = 0
contains
  subroutine foo create (this, ...)
    type(foo) :: this
    if (.not. allocated(this%stuff) then
       : ! allocate and initialize
      foo count = foo count + 1
                                         must explicitly invoke
    end if
  end subroutine foo create
  subroutine foo destroy (this)
    foo count = foo count - 1
  end subroutine
end module mod foo
```

module is encapsulation unit

C++ uses a static member variable

```
class Foo {
  public:
    Foo() : len_(0),stuff_(NULL) {};
    Foo(int, float *);
    :
  protected:
    static int count;
    int len_;
  private:
    float *stuff_; };
```

```
#include "Foo.h"
int Foo::count = 0;

Foo::Foo() {
   count += 1;
}
// same with all other constr.
// decrement in destructor
```

class is encapsulation unit

Updates to a shared entity

... and don't forget to switch on OpenMP everywhere!



Example:

execute object creation in parallel region

```
type(foo) :: obj
:
! obj not created yet here
!$omp parallel private(obj, ...)
call foo_create(obj, ...)
: ! do computations
call foo_destroy(obj)
!$omp end parallel
! obj undefined
```

- beware definition status
- updates on foo_count are not thread-safe
- → inconsistencies / wrong values

Fix: use a named critical

```
subroutine foo_create(this, ...)
:
!$omp critical (c_count)
    foo_count = foo_count + 1
!$omp end critical
    end subroutine foo_create

subroutine foo_destroy(this)
:
!$omp critical (c_count)
    foo_count = foo_count - 1
!$omp end critical
    end subroutine foo_destroy
```

 imagine foo_count is public → need an efficient tool to identify any problem

The environment problem: setting the stage



Calculation of

$$I = \int_{a}^{b} f(x, p) dx$$

where

- f(x,p) is a real-valued function of a real variable x and a variable p of some undetermined type
- a, b are real values
- Tasks to be done:
 - procedure with algorithm for establishing the integral \rightarrow depends on the properties of f(x,p) (does it have singularities? etc.)

$$I \approx \sum_{i=1}^{n} w_i f(x_i, p)$$

- function that evaluates f(x,p)
- Case study provides a simple example of very common programming tasks with similar structure in scientific computing.

Using a canned routine: D01AHF (Patterson algorithm in NAG library)



Interface:

requested precision

```
double precision FUNCTION D01AHF (A, B, EPSR, NPTS, RELERR, F, NLIMIT, IFAIL)
   INTEGER :: NPTS, NLIMIT, IFAIL
   double precision :: A, B, EPSR, RELERR, F
   EXTERNAL :: F
```

uses a function argument

```
double precision FUNCTION F (X)
    double precision :: X
```

(user-provided function)

Invocation:

```
define a, b

res = d01ahf(a, b, 1.0e-11, &
    npts, relerr, my_fun, -1, is)
```

Mass-production of integrals

may want to parallelize

```
!$omp parallel do
do i=istart, iend
   :! prepare
   res(i) = d01ahf(..., my_fun, ...)
end do
!$omp end parallel do
```

 need to check library documentation: thread-safeness of d01ahf

Mismatch of user procedure implementation



User function may look like this:

```
subroutine user_proc(x, n, a, result)
  real(dk), intent(in) :: x, a
  integer, intent(in) :: n
  real(dk), intent(out) :: result
end subroutine
```

- parameter "p" is actually the tuple $(n, a) \rightarrow no$ language mechanism available for this
- or like this

```
real(dk) function user_fun(x, p)
    real(dk), intent(in) :: x
    type(p_type), intent(in) :: p
end function
Compiler would accept
    this one due to the
    implicit interface for it,
    but it is likely to bomb at run-time
```

Neither can be used as an actual argument in an invocation of dolahf

Solution 1: Wrapper with global variables



Usage:

```
use mod_user_fun

par = ...; n = ...
res = d01ahf(..., arg_fun, ...)

supply values
for global variables
```

Disadvantages of Solution 1



Additional function call overhead

 is usually not a big issue (nowaday's implementations are quite efficient, especially if no stack-resident variables must be created).

Solution not thread-safe (even if d01ahf itself is)

expect differing values for par and n in concurrent calls:

```
!$omp parallel do
do i=istart, iend
  par = ...; n = ...
  res(i) = d01ahf(..., arg_fun, ...)
end do
!$omp end parallel do

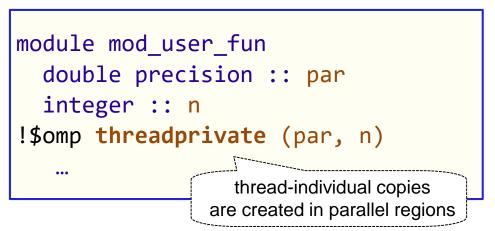
join
```


 results in unsynchronized access to the shared variables par and n from different threads → race condition → does not conform to the OpenMP standard → wrong results

Making Solution 1 thread-safe



Threadprivate storage





Usage may require additional care as well

```
par = ...
!$omp parallel do copyin(par)
do i = istart, iend
    n = ...
    ... = d01ahf(..., arg_fun, ...)
    if (...) par = ...
end do
!$omp end parallel do

broadcast from master copy
needed for par


A bit cumbersome:
non-local programming
style required
```

Solution 2: Reverse communication



Change design of integration interface:

- instead of a function interface, provider requests a function value
- provider provides an argument for evaluation, and an exit condition



Solution 2: Typical example interface



Uses two routines:

```
subroutine initialize_integration(a, b, eps, x)
    real(dk), intent(in) :: a, b, eps
    real(dk), intent(out) :: x
    means?
end subroutine
subroutine integrate(fval, x, result, stat)
    real(dk), intent(in) :: fval
    real(dk), intent(out) :: x
    real(dk), intent(inout) :: result
    integer, intent(out) :: stat
end subroutine
do you remember
    what "INTENT"
    means?

shall not be modified by caller
    while calculation iterates
```

- first is called once to initialize an integration process
- second will be called repeatedly, asking the client to perform further function evaluations
- final result may be taken once stat has the value stat_continue

Solution 2: Using the reverse communication interface



```
program integrate
  real(dk), parameter :: a = 0.0_{dk}, b = 1.0_{dk}, eps = 1.0e-6_dk
  real(dk) :: x, result, fval, par
  integer :: n, stat
  n = ...; par = ...
  call initialize_integration(a, b, eps, x)
  do
    call user_proc(x, n, par, fval)
    call integrate(fval, x, result, stat)
    if (stat /= stat_continue) exit
  end do
 write(*, '(''Result: '',E13.5,'' Status: '',I0)') result, stat
contains
  subroutine user proc( ... )
  end subroutine user_proc
end program
```

- avoids the need for interface adaptation and global variables
- some possible issues will be discussed in an exercise



Taking Solution 2 a step further



Disadvantage:

- iteration routine completes execution while algorithm still executes
- this may cause a big memory allocation/deallocation overhead if it uses many (large) stack (or heap) variables with local scope
- Note: giving such variables the SAVE attribute causes the iteration routine to lose thread-safeness

Concept of "coroutine"

- type of procedure that can interrupt execution without deleting its local variables
- co-routine may return (i.e. complete execution), or suspend
- invocation may call, or resume the coroutine

(implies rules about invocation sequence)

- no language-level support for this exists in Fortran
- however, it can be emulated using OpenMP



Coroutine emulation via OpenMP tasking



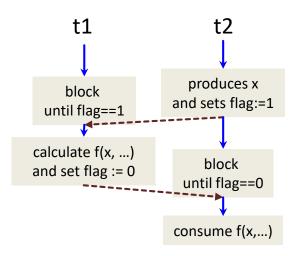
Separate tasks are started for

- supplier, and for
- consumer of function values

```
n = ...; par = ...; a = ...; b = ...; eps = ...
 flag = flag need iter
!$omp parallel num_threads(2) proc_bind(master)
!$omp single
                     task t1
!$omp task ...
    do
      call user_proc(x, n, par, fval)
    end do
!$omp end task
!$omp task t2
    call integrate_c(a, b, eps, fval, x, &
                      result, flag)
!$omp end task
!$omp end single
                         continues executing until
!$omp end parallel
                        the algorithm has completed
```

Explicit synchronization needed

- between supplier and consumer
- functional (vs. performance) threading
- involved objects: x, fval
- use an integer flag for synchronization





Synchronization code



Look at task block "t1" from previous slide in more detail:

```
!$omp task private(flag local)
!$omp taskyield
   iter: do
      spin: do
!$omp atomic read
         flag local = flag
         if (flag_local == flag_need_fval) exit spin
         if (flag local > 1) exit iter
     end do spin
!$omp flush(x)
     call user proc(x, n, par, fval)
!$omp flush (fval)
!$omp atomic write
     flag = flag_need_iter
!$omp taskyield
   end do iter
!$omp end task
```

- A mirror image of this is done inside integrate_c()
- Grey area with respect to Fortran conformance (aliasing rules)

the TARGET attribute might help



Dynamic memory and object-based design

Recapitulation: dynamic objects



Add a suitable attribute to an entity:

```
initial state is "unallocated"
                                                        initial state is "unassociated"
 real, allocatable :: x(:)
                                           real, pointer :: p(:) => null()
               deferred shape
Typical life cycle management:
                                                  non-default lower bounds are possible
                                                  (use LBOUND and UBOUND intrinsics)
                    allocate(x(2:n), p(3), stat=my_status)
              create
 use of
                    x(:) = ...
                                            definitions and references
                    p(:) = ...
 heap
memory
                    deallocate(x, p)
              destrov
```

Status checking:

(hints at semantic differences!)

```
if (allocated(x)) then; ... if (associated(p)) then; ...
```


ALLOCATABLE vs. POINTER




- An allocated allocatable entity
 - is an object in its own right
 - becomes auto-deallocated once going out of scope
- An associated pointer entity
 - is an alias for another object, its target
- real, target :: tg(3) = 0.0

except if object has the SAVE attribute e.g., because it is global

all definitions and references are to the target




undefined (third) state should be avoided



Implications of POINTER aliasing





Multiple pointers may point to the same target

Avoid dangling pointers

```
p2 is associated with all of the target.
p1 and p3 become undefined deallocate(p2)
nullify(p1, p3)
```

Not permitted: deallocation of allocatable target via a pointer

```
real, allocatable, target :: t(:) all real, pointer :: p(:) dea
```

Features added in





Allocatable entities

Scalars permitted:

```
real, allocatable :: s
```

 LHS auto-(re)allocation on assignment:

```
x = p(2:m-2)
```

• The MOVE_ALLOC intrinsic:

```
call MOVE_ALLOC(from, to)
```

Pointer entities

rank changing "=>":

```
real, target :: m(n)
real, pointer :: p(:,:)
p(1:k1,1:k2) => m

rank of target must be 1
```

bounds changing "=>":

```
p(4:) => m

bounds remapped via
lower bounds spec
```

Deferred-length strings:

pointer also permitted, but subsequent use is then different




Not a Fortran term Container types (1)



Allocatable type components

```
type :: polynomial
  private
  real, allocatable :: f(:)
end type
             default (initial) value is
                 not allocated
```


a "value" container



POINTER type components

```
type :: sparse
  private
  integer :: index
                        default value is
  real :: value
                        disassociated
  type (sparse), &
     pointer :: next => null()
end type
```

a "reference" container



example type is selfreferential → "linked list"

Container types (2): Object declaration and assignment semantics



Allocatable type components

 assignment statement is equivalent to

"deep copy"

POINTER type components

```
type(sparse) :: s1, s2
: define s1
s2 = s1
```

assignment statement is equivalent to

"shallow copy"

Container types (3): Structure constructor



Allocatable type components

```
type(polynomial) :: p1
p1 = polynomial([1.0, 2.0])
```

 dynamically allocates p1%f to become a size 2 array with elements 1.0 and 2.0

When object becomes undefined

 allocatable components are automatically deallocated

usually will not happen for POINTER components

POINTER type components

```
type(sparse) :: s1
type(sparse), target :: t1
type(sparse), &
    parameter :: t2 = ...
s1 = sparse( 3, 1.0, null() )
```

• alternative:

```
s1 = sparse( 3, 1.0, t1 )
```

not permitted:

```
s1 = sparse( 3, 1.0, t2 )
```

a constant cannot be a target

→ e.g., overload constructor to avoid this situation (create argument copy)

Container types (4): Storage layout




Irregularity:

- each array element might have a component of different length
- or an array element might be unallocated (or disassociated)

```
type(polynomial) :: p_arr(4)

p_arr(1) = polynomial( [1.0] )
p_arr(3) = polynomial( [1.0, 2.0] )
p_arr(4) = polynomial( [1.0, 2.0, 3.1, -2.1] )
```



- Applies for both allocatable and POINTER components
 - a subobject designator like p_arr(:)%f(2) is not permitted

Allocatable and POINTER dummy arguments



(explicit interface required)

Allocatable dummy argument

useful for implementation of "factory procedures" (e.g. by reading data from a file)

POINTER dummy argument

example: handling of a "reference container"

Actual argument must have matching attribute

an exception to this will be mentioned

INTENT semantics for dynamic objects



specified intent	allocatable dummy object	pointer dummy object	
in	procedure must not modify argument or change its allocation status	procedure must not change association status of object	
out	argument becomes deallocated on entry auto-deallocation of simulation_field on previous slide!	pointer becomes undefined on entry	
inout	retains allocation and definition status on entry	retains association and definition status on entry	

"Becoming undefined" for objects of derived type:

- type components become undefined if they are not default initialized
- otherwise they get the default value from the type definition
- allocatable type components become deallocated

INTENT(OUT) and default initialized types



Suppose that a derived type person has default initialization:

```
type :: person
  character(len=32) :: name = 'no_one'
  integer :: age = 0
end type
```

then, after invocation of

```
subroutine modify_person(this)
  type(person), intent(out) :: this
  :
  this%name = 'Dietrich'
  ! this%age is not defined
end subroutine
```

the actual argument would have the value person('Dietrich',0), i.e. components not defined inside the subprogram will be set to their default value

Quiz: what happens with a POINTER component in this situation?

Bounds of deferred-shape objects



- Bounds are preserved across procedure invocations and pointer assignments
 - Example:

```
real, pointer :: my_item(:) => null
type(container), intent(inout) :: my_container(ndim)
allocate(my_item(-3:8))
call add_reference(my_container(j), my_item)
```

What arrives inside **add_reference**?

- this is different from assumed-shape, where bounds are remapped
- it applies for both POINTER and ALLOCATABLE dummy objects

CONTIGUOUS pointers



The CONTIGUOUS attribute can be specified for pointers

- (and also for assumed-shape arrays)
- difference to assumed-shape: programmer is responsible for guaranteeing the contiguity of the target in a pointer assignment

Example:

also illustrates rank changing:

```
real, pointer, contiguous :: matrix(:,:)
:
allocate(storage(n*n))
matrix(lb:ub,lb:ub) => storage
```

matrix can be declared contiguous because whole allocated array storage is contiguous

if contiguity of target is not known, check via intrinsic:

Allocatable function results



(explicit interface required)

Scenario:

- size of function result cannot be determined at invocation
- example: remove duplicates from array

Possible invocations:

efficient (uses auto-allocation on assignment):

```
integer, allocatable :: res(:)
res = deduplicate(array)
```

 less efficient (two function calls needed):

 function result is auto-deallocated after completion of invocation

POINTER function results

(explicit interface required)



POINTER attribute

for a function result is permitted,



it is more difficult to handle on both the provider and the client side (need to avoid dangling pointers and potential memory leaks)

Example: filtering a list

```
function next uppertr(s, i) result(r)
  type(sparse), target, intent(in) :: s
  integer, intent(in) :: i
  type(sparse), pointer :: r
  r \Rightarrow s
                             code to identify first
  do while (...)
                            entry with index >= i
    r => r%next ----
                              pointer assignment
                              to existing TARGET
  end do
  if (...) r => null()
end function next uppertr
```

invocation:

```
type(sparse), target :: trm(nd)
type(sparse), pointer :: entry
       set up my_matrix (linked list)
do i=1, nd
  entry => trm(i)
  do while ( associated(entry) )
      do work on entry
    entry => next_uppertr(entry,i)
end do
```

- note the pointer assignment
- it is essential for implementing correct semantics and sometimes also to avoid memory leaks

based on earlier opaque type definition of sparse

Opinionated recommendations



Dynamic entities should be used, but sparingly and systematically

- performance impact, avoid fragmentation of memory → allocate all needed storage at the beginning, and deallocate at the end of your program; keep allocations and deallocations properly ordered.
- If possible, ALLOCATABLE entities should be used rather than POINTER entities
 - avoid memory management issues (dangling pointers and leaks)
 - especially avoid using functions with pointer result
- A few scenarios where pointers may not be avoidable:
 - information structures → program these in an encapsulated manner: user of the facilities should not see a pointer at all, and should not need to declare entities targets.
 - subobject referencing (arrays and derived types) → performance impact!

Recapitulation: Generic procedures



Named interfaces

```
interface generic_name
  procedure :: specific_1
  procedure :: specific_2
  ...
end interface
```

- signatures of any two specifics must be sufficiently different (compile-time resolution)
- Potential restrictions on signatures of specific procedures

Operator overloading or definition

```
interface operator (+)
  procedure :: specific_1
  procedure :: specific_2
  ...
end interface
```

```
interface operator (.user_op.)
  procedure :: specific_1
  procedure :: specific_2
  ...
end interface
```

- operators: functions with two arguments (one for unary operations)
- assignment: subroutine with two arguments
- overloaded structure constructor: function with type name as result
- user-defined derived type I/O (treated on day 2)

Generalizing generic interface blocks



can be replaced by

```
interface foo_generic
  module procedure foo_1
  module procedure foo_2
end interface
```

with generalized functionality:

```
interface foo_generic
  procedure foo_1
  procedure foo_2
end interface
```

Referenced procedures can be

- external procedures
- dummy procedures
- procedure pointers

Example:

```
interface foo_gen
! provide explicit interface
! for external procedure
  subroutine foo(x,n)
    real, intent(out) :: x
    integer, intent(in) :: n
  end subroutine foo
end interface
interface bar_gen
  procedure foo
end interface
```

- is valid in
- . .
- is non-conforming if a

```
module procedure
```

statement is used

Case study - sparse matrix operations



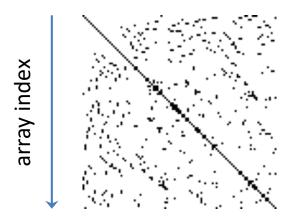
Represent sparse matrix

type(sparse), allocatable :: sa(:)

- sa(i) is the i-th row of the matrix
- sa(i)%value is the non-zero value of the sa(i)%index column element
- sa(i)%next is associated with the next non-zero entry

Occupancy graph

 non-zero elements represented by black dots



- Creating, copying and operations of such objects
 - topics for the next slides and the exercises

Overloading the structure constructor (F03)





Rationales:

- default structure constructor not generally usable due to encapsulation of type components
- default structure constructor cannot by itself set up complete list or array structures
- input data characteristics may not match requirements of default constructor

```
module mod sparse
  : ! previous type definition for sparse
  interface sparse generic has same name as the type
    procedure :: create sparse
                                         more than one specific is possible
  end interface
contains
                                 must be a function with scalar result
  function create_sparse(colidx, values) result(r)
    integer, intent(in) :: ncol(:), colidx(:)
    real, intent(in) :: values(:)
    type(sparse) :: r
                                     implementation dynamically allocates
                                          the linked list for each row
  end function
end module mod sparse
```

Notes on overloading the structure constructor



- If a specific overloading function has the same argument characteristics as the default structure constructor, the latter becomes unavailable
 - advantage: for opaque types, object creation can also be done in use association contexts
 - disadvantage: it is impossible to use the overload in constant expressions

Of course, a specific may have a wildly different interface, corresponding to the desired path of creation for the object (e.g., reading it in from a file)

Applying default assignment properties



For the overloaded constructor, ...

- ... would work fine if A(i) was not previously established)
- However, for a "regular" assignment,

```
type(sparse), allocatable :: A(:), B(:)
:
    A(i) = sparse(colidx, values)
:
B = A
RHS persis
```

assignment, but not the allocated component memory

A(i) = sparse(...)

function result is discarded after

anonymous target of next

RHS persists after the assignment

- B effectively is not an object in its own right, but (except for the first array element in each row) links into A.
- Also, default assignment is unavailable between objects of different derived types

Overloading the assignment operator



Uses a restricted named interface:

```
module mod sparse
     ! type definition of sparse
  interface assignment(=)
    procedure assign_sparse
 end interface
                  exactly two arguments
contains
  subroutine assign_sparse(res, src)
    type(sparse), intent(out) :: res
    type(sparse), intent(in) :: src
      implement a deep copy
 end subroutine
end module
```

create a clone of the RHS

Further rules:

- first argument: intent(out)or intent(inout)
- second argument: intent(in)
- assignment cannot be overloaded for intrinsic types
- overload usually wins out vs. intrinsic assignment.
 Exception: implicitly assigned aggregating type's components → aggregating type must also overload the assignment


Quiz: what might be missing in the procedure definition?

Overloaded assignment of function results: Dealing with POINTER-related memory leaks



Scenario:

 RHS may be an (overloaded) constructor or some other function value (e.g. an expression involving a defined operator)



Fos Therapy:

- add a finalizer to type definition
- references a module procedure with a restricted interface (usually, a single scalar argument of the type to be finalized)

Finalizing procedure implementation



```
applicability to array objects

elemental recursive subroutine finalize_sparse(this) type(sparse), intent(inout) :: this if (associated(this%next)) then deallocate(this%next) assumes that all targets end if
```

Implicit execution of finalizer:

end subroutine

- when object becomes undefined (e.g., goes out of scope),
- is deallocated,
- is passed to an intent(out) dummy argument, or
- appears on the left hand side of an intrinsic assignment

```
Quiz: what happens in the assignment

A(i) = sparse(...)

if a finalizer is defined, but the assignment is not overloaded?
```



allocated

Notes on finalizers



- Feature with significant performance impact
 - potentially large numbers of invocations:
 array elements, list members
 - finalizer invoked twice in assignments with a function value as RHS
- Finalizers of types with pointer components:
 - may need to consider reference counting to avoid undefined pointers
- Non-allocatable variables in main program
 - have the implicit SAVE attribute → are not finalized
- Further comments on finalizers will be made on day 4

following now: Exercise session 2

Recall aliasing of dummy arguments



Definition

- access to object (or subobject) via a name other than the argument's name:
- (sub)object of actual argument is associated with another actual argument (or part of it)
- 2. actual argument is (part of) a global variable which is accessed by name
- 3. actual variable (or part of it) can be accessed by host association in the executed procedure

Simplest example:

illustrates item 1

```
subroutine foo(a, b)
  real, intent(inout) :: a
  real, intent(in) :: b
  a = 2 * b
  ... = b
  a = a + b
end subroutine
```

some invocations:

```
real :: x, y
x = 2.0; y = x
call foo(x, x)
! aliased - non-conforming
call foo(x, y)
! not aliased - x is 6.0
call foo(x, (x))
! not aliased - x is 6.0
```

Aliasing restrictions



Restriction 1:

 if (a subobject of) the argument is defined in the subprogram, it may not be referenced or defined by any entity aliased to that argument

Notes:

- this restriction renders the first call illegal
- but aliasing is not generally disallowed
- exceptions to this rule will be discussed later

Intent:

enable performance optimizations by statement reordering and/or register use

 avoid ambiguities in assignments to dummy arguments

Restriction 2:

- changes of allocation or association status of (part of) a dummy argument may only be performed via this argument.
- subsequent to such a change, any references or definitions of the object may be only via this argument

Note:

 deals with expected semantics of handling descriptor passing for allocatable or POINTER objects (usually copy-in/out)

Diagnosis of aliasing:

Requires inspection of procedure implementation as well as its invocation

- invocation for whether aliasing occurs
- implementation for whether the restrictions are violated

Violation scenario for Restriction 2



A non-conforming variation on the factory method seen earlier

```
program simulation
                                                           copy in/out of
            implicit none
                                                         descriptor for field
            real, allocatable :: field(:,:,:)
                                                           might be done
            call read simulation data(field, 'my f.dat')
                                                                      argument
                                                                     association
          contains
            subroutine read simulation data(simulation field, file name)
field is host
               real, allocatable, intent(out) :: simulation_field(:,:,:)
 associated
               character(len=*), intent(in) :: file name
               allocate(simulation_field(...)) ! and fill in values
               if ( .not. allocated(field) ) &
                    allocate(field(...)) | and possibly give it values
            end subroutine read_simulation_data
                                                                   this statement
                                                                   may well get
          end program
                                                                executed and succeed
```

- and after return from the procedure further bad effects will occur
- interdiction against this applies for ALLOCATABLE and POINTER objects

Exceptions to Restriction 1: POINTER dummy arguments



By definition, pointers implement aliasing to their target

 hence, for a dummy argument with the POINTER attribute restriction 1 does not hold

(a pointer can be regarded as an orphaned dummy argument)

Note that:

- the aliasing property implies that the POINTER attribute has a negative performance impact;
- the TARGET attribute on an object indicates to the compiler that pointers may be associated with the object or part of it. Optimization may depend on whether this currently is the case.

Example for permitted aliasing



The following program is conforming

```
program aliasing 1
                                                     both actual and dummy
            implicit none
                                                      argument have the
            real, pointer :: p(:)
                                                      POINTER attribute
            allocate (p(-10:10))
            call modify_ptr(p)
            deallocate(p)
                                                        p is argument
p is host
          contains
                                                       associated with x
associated
            subroutine modify ptr(x)
               real, pointer, intent(in) :: x(:)
               integer :: i
                                                                compiler cannot
               do i = lbound(p,1) + 1, ubound(p,1)
                                                               vectorize this code
                 p(i) = p(i) + x(i - 1)
                                                          (effectively, a flow dependency)
               end do
            end subroutine modify_ptr
          end program
```

note the explicit interface

Exceptions to Restriction 1: Dummy arguments with the TARGET attribute



Restriction 1 is lifted under the following additional conditions:

1. Dummy argument

- is not intent(in) or value
- is a scalar or an assumedshape array

2. Actual argument

- also has the TARGET attribute
- is not an array section with a vector subscript

These conditions

suppress copy-in/out

and

 preserve pointer association across the interface (if the ultimate actual argument has the POINTER attribute, or a global pointer is associated with the dummy argument)

Second example for permitted aliasing



The following program is conforming


```
program aliasing_2
            implicit none
            real, pointer :: p(:)
            allocate (p(-10:10)
            call modify tgt(p)
                                       target of p is argument
            deallocate(p)
                                          associated with x
          contains
p is host
            subroutine modify_tgt(x)
associated
              real, target, intent(inout) :: x(:)
              integer :: i, ip
              do i = 2, size(x, 1)
                                                              compiler cannot
                  ip = i + lbound(p,1) - 1
                                                             vectorize this code
                                                        (effectively, a flow dependency)
                  p(ip) = p(ip) + x(i-1)
              end do
            end subroutine modify tgt
          end program
```

note the explicit interface

Example for forbidden aliasing



The following program is non-conforming



• inside the procedure, associated(p, x) may return false or true

Interfaces: temporarily acquiring or losing the TARGET attribute



Generally, an explicit interface is required

 for having the TARGET attribute on a dummy argument

Case 1:

dummy argument has the attribute, but actual does not

- then, any pointer associated with the target during execution of the subprogram becomes undefined at its completion
- association is with the dummy argument only, not the actual argument

Case 2:

actual argument has the attribute, but dummy argument does not

- then, pointer associations with the actual argument are not affected
- but association with dummy argument is undefined

Case 3: (example on previous slide)

 if the additional conditions for aliasing permission are not fulfilled, pointer association is not guaranteed to be preserved across the invocation / completion

Fortran POINTERs: Handling the argument association



Actual Argument	Dummy Argument object pointer target		target	
object	usually by reference, may need copy-in/ copy-out (efficiency), no-alias assumption	not allowed	pointer assoc. with dummy argument becomes undefined on return	
pointer	must be associated, dereference to target, may need copy-in/ copy-out (efficiency)	same rank, associa- tion status passed, beware invalid target (upon return)	is preserved .	
target	usually by reference, copy-in/copy-out allowed, no-alias assumption	permitted in 🕬 if dummy is intent(in)	copy-in/copy-out not allowed for scalar and assumed shape array dummies	
		explicit interface required		

An alternative aliasing mechanism





Alternative: association block

 combine aliasing with a block construct to avoid pointerrelated performance problems

Association syntax fragment:

(<associate name> => <selector>)

 allows to use the associate name as an alias for the selector inside the subsequent block

Very useful for

- heavily reused complex expressions (especially function values)
- references into deeply nested types

Selector:

- may be a variable → associate name is definable
- may be an expression → is pre-evaluated before aliasing to associate name, which may not be assigned to

Inherited properties:

- type, array rank and shape, polymorphism (discussed later)
- asynchronous, target and volatile attributes

Not inherited:

pointer, allocatable and optional attributes

Block construct ASSOCIATE




Example:

 given the type definitions and object declaration:

```
type :: vec_3d
  real :: x, y, z
end type
type :: system
  type(vec_3d) :: vec
end type
type :: all
  type(system) :: sys
  real :: q(3)
end type
type(all) :: w
```

 the following block construct can be established



Notes:

- more than one selector can be aliased for a single block
- the associate is auto-typed (an existing declaration in surrounding scope becomes unavailable)
- writing this out in full would be very lengthy and much less readable

Recommendations for library design (1)



Library not a static entity

- may want to add to functionality
- may need to fix bugs / design problems
- Open/Closed principle (OCP)

Any software entity should be

- open for extension
- closed against modification

B. Meyer (1988)

- client using the library should not run into trouble
- at minimum, client source code should build against updated library and execute as prior to the update

higher level of "closed":
 binary compatibility - either replacement of shared libraries or relinking is sufficient

Assumption for the following discussion:

 all library code is implemented in form of modules

(modules have improved support for many aspects of software engineering)

Recommendations for library design (2)



Changes to implementations

- typically bug fixes in bodies of module procedures
- interfaces (procedure signature) unchanged

Consequences:

- theoretically, relinking against the library should be sufficient
- if compilation of client code is performed, recompilation of all units directly or indirectly depending on the changed module must be done ("cascade")

Changes to existing interfaces

- normally forbidden
- may be able to circumvent incompatibility via introduction of a generic

```
one specific may have the generic name module procedure my_sub module procedure my_corrected_sub end interface in specification part
```

(or an optional argument)

add my_corrected_sub() as a module procedure

Consequence:

 need to recompile all dependent clients and relink

Recommendations for library design (3)



Derived types

- keep type components private ("information hiding")
- exposed type components cannot be changed → would typically render client code unworking, therefore violates the OCP

Global data

- Declaration (name, most attributes) cannot be changed if public (for the same reason)
- Consequence of changes on private type components or global data
 - need to recompile all dependent clients and relink

following now: Exercise session 3



Object-oriented programming (I)

Type extension and polymorphism

Characterization



Terminology

- terms and their meaning vary between languages → danger of misunderstandings
- will use Fortran-specific nomenclature (some commonly used terms may appear)

Aims of OO paradigm: improvements in

- re-using of existing software infrastructure
- abstraction
- moving from procedural to data-centric programming
- reducing software development effort, improving productivity

Indiscriminate usage of OO however may be (very) counterproductive

identify "software patterns" which have proven useful

Scope of OO within Fortran



- Fortran 95 supported object-based programming
- Today's Fortran supports object-oriented programming
 - type extension and polymorphism (single inheritance)
 - type-bound and object-bound procedures, finalizers and type-bound generics
 - extensions to the interface concept


Specific intentions of Fortran object model:

- backward compatibility with Fortran 95
- allow extensive correctness and consistency checking by the compiler
- module remains the unit of encapsulation, but encapsulation becomes more fine-grained
- design based on Simula object model

Type extension (1): Defining an extension



Type definitions



- idea: re-use date definition
- datetime a specialization (or Prerequisite: subclass) of date
- date more general than datetime

Fortran type extension

```
type :: date
  private
  integer :: year = 0
  integer :: mon = 0, day = 0
end type
type, extends(date) :: datetime
  private
  integer :: hr = 0, min = 0, &
             sec = 0
end type
```

single inheritance only

- parent type must be **extensible**
- i.e., be a derived type that has neither the SEQUENCE nor the BIND(C) attribute

Type extension (2): Declaring an object of extended type



If type definition is public

 an object of the extended type can be declared in the host, or in a program unit which use associates the defining module

```
use mod_date
:
type(datetime) :: o_dt
```

Accessing component data

inherited components:

```
o_dt%day o_dt%mon o_dt%yr
```

additional components

```
o_dt%hr o_dt%min o_dt%sec
```

Parent component

```
o_dt % date
```

- is an object of parent type
- a subobject of o_dt
- recursive references possible:o_dt % date % day
- parent components are themselves inherited to further extensions

Note:


 encapsulation may limit accessibility for all component variants

Type extension (3): General form of inheritance tree



A directed acyclical graph (DAG)

this is a consequence of supporting single inheritance only



Variants:

- flat inheritance tree (typically only one level)
 - base type is provided, which everyone else extends
 - very often with an abstract type (discussed later) as base type
- deep inheritance tree
 - requires care with design (which procedures are provided?) and further extension
 - requires thorough documentation

Type extension (4): Further notes



Extension can have zero additional components

use for type differentiation:

```
type, extends(date) :: mydate
end type
type(mydate) :: o_mydate
```

- o_mydate cannot be used in places where an object of type(date) is required
- or to define type-bound procedures (discussed later) not available to parent type

Type parameters are also inherited

 see later slide for more details

Inheritance and scoping:

 cannot have a new type component or type parameter in an extension with the same name as an inherited one

(name space of class 2 identifiers)


Type extension (5): Component accessibility issues



Example: A type extension defined via use association

Inheritance of accessibility:

- o_dth has six inherited
 private components and one
 public one
- supports mixed accessibility of type components!



- Technical Problem (TP1) for opaque types:
 - cannot use the structure constructor for datetime_hires
 - reason: it is only available outside the host of mod_date, hence privateness applies
 - one solution: overload structure constructor

Explicit syntax for mixed component access





Example: a partially opaque derived type

```
module mod_person
  type :: person
  private
  character(len=strmx) :: name
  integer :: age
  character(len=tmx), public :: location
  end type
  :! module procedures are not shown
  end module

  design decision: location
  is not encapsulated. Why?
```

any program unit may modify the %location component:

```
use mod_person, only : person
type(person) :: p
: ! initialize p via an accessor defined in mod_person
p%location = 'room E.2.24' ! update location
```

Type extension (6): Structure constructor



Using keywords

 example: inside the host of mod_date, one can have

- → change component order
- rules are as for procedure keyword arguments
- e.g., once keyword use starts, it must be continued for all remaining components

Using parent component construction

 example: inside the host of mod_date, one can have

keyword notation required!

General restriction:

 it is not allowed to write overlapping definitions, or definitions that result in an incomplete object

Further structure constructor features in (F03)





Omitting components in the structure constructor

- this omission is only allowed for components that are default-initialized in the type definition
- example: in any program unit, one can have

```
use mod ext
type (datetime_hires) :: o_hires
o hires = datetime hires(msec=711)
```

because all other components will receive their defaultinitialized value

- also applies to POINTER and ALLOCATABLE components (further details on day 3)
- sometimes, this alleviates the TP1 from some slides earlier

Polymorphism (1): Polymorphic objects





Declaration with CLASS:

class(date), ... :: o_poly_dt

possible additional attributes

- declared type is date
- dynamic type may vary at run time: may be declared type and all its (known) extensions (type compatibility)

Data item can be

1. dummy data object

interface polymorphism

2. pointer or allocatable variable

data polymorphism → a new kind of dynamic memory

3. both of the above

loosening of strict F95 typing rules

 direct access (i.e., references and definitions) only possible to components of declared type (compile-time: compiler lacks knowledge, run-time: semantic problem and performance issues)

invalid even if dynamic type of o poly dt is datetime

Polymorphism (2): Interface polymorphism



Example:

- increment date object by a given number of days
- Inheritance mechanism: actual argument ...
 - ... can be of declared type of dummy or an extension:

```
type(date) :: o1
type(datetime) :: o2
: ! initialize both objects
call inc_date(o1,2._rk)
call inc_date(o2,2._rk)
```

 ... can be polymorphic or nonpolymorphic

```
subroutine inc_date(this, days)
  class(date), intent(inout) :: this
  real(rk), intent(in) :: days
  :! implementation → exercise
end subroutine
```

```
could replace "type(...)" by
"class(...)" for both objects
(an additional attribute may be needed)
```

Argument association:

 dynamic type of actual argument is assumed by the dummy argument

Polymorphism (3): Interface polymorphism cont'd

assume



Example continued:

account for fraction of a day when incrementing a datetime object

Restriction on use:

cannot take objects of declared type date as actual argument:

```
class(date) :: o1
                                dummy
class(datetime) :: o2
                               arguments
  ! initialize both objects
call inc_datetime(o1,.03_rk)
                                  invalid invocation -
call inc_datetime(o2,.03_
                                    will not compile
                                   (this also applies if o1
                                   is of non-polymorphic
                                      type(date))
```

```
subroutine inc_datetime(this, days)
  class(datetime), &
             intent(inout) :: this
  real(rk), intent(in) :: days
    ! implementation → exercise
end subroutine
```

reason: if o1 has dynamic type date, then no sec component exists that can be incremented

Fortran term:

dummy argument must be type compatible with actual argument

(note that type compatibility, in general, is not a symmetric relation)

Polymorphism (4): Data polymorphism / dynamic objects



Declaration:

- unallocated / disassociated entities: dynamic type is equal to declared type
- usual difference in semantics (e.g., auto-deallocation for allocatables)

Producing valid entities:

 typed allocation to base type or an extension

```
allocate(datetime :: ad, cd)
becomes dynamic type

allocate(date :: bd(5))

could omit since equal to base type
```

pointer association

Polymorphism (5): Arrays



A polymorphic object may be an array

```
class(date) :: ar_d(:)
```

here: assumed-shape

(Note: using assumed-size or explicit-shape is usually not a good idea)

but type information applies for all array elements

all array elements have the same dynamic type

For per-element type variation:

 define an array of suitably defined derived type:

```
type :: date_container
  class(date), allocatable :: p
end type

type(date_container) :: arr(10)
```

arr(1)%p can have a dynamic type different from that of arr(2)%p

Polymorphism (6): Further allocation mechanisms



object ad: declared two slides earlier

Sourced allocation

produce a **clone** of a variable or expression

```
class(datetime) :: src
 ! define src
allocate(ad, source=src)
```

- allocated variable (ad) must be type compatible with source
- source can, but need not be polymorphic
- definition of dynamic type of source may be inaccessible in the executing program unit (!)
- usual semantics: deep copy for allocatable components, shallow copy for pointer components

Sourced allocation of arrays

array bounds are also transferred in sourced allocation

Molded allocation (FOS)



allocate an entity with the same shape, type and type parameters as mold

```
class(datetime) :: b
allocate(ad, mold=b)
```

- mold need not have a defined value (no data are transferred)
- otherwise, comparable rules as for sourced allocation

Polymorphism (7): Type resolution



Example scenario:

 a routine is needed that writes a complete object of class(date) to a file irrespective of its dynamic type

```
subroutine write_date(this, fname)
  class(date), intent(in) :: this
  character(len=*) :: fname
  :! open file fname on unit
    :! see inset right
end subroutine
```

Problem:

 how can extended type components be accessed within write date?

New block construct:

```
must be polymorphic
select type (this)
type is (date)
  write(unit,fmt='("date")')
  write(unit,...) this%day,...
type is (datetime)
  write(unit,fmt='("datetime")')
  write(unit,...) ...,this%hr,...
        inside this type guard block:
        this is nonpolymorphic

    type of this is datetime

  ! further type guards for
  ! other extensions
class default
  stop 'Type not recognized'
end select
                  fall-through block:

    this is polymorphic

    typically used
```

Polymorphism (8): Semantics and rules for SELECT TYPE



Execution sequence:

- at most one block is executed
- selection of block:
- 1.find **type guard** ("type is") that exactly matches the dynamic type
- 2.if none exists, select **class guard** ("class is") which most closely matches dynamic type and is still type compatible
 - → at most one such guard exists
- 3.if none exists, execute block of class default (if it exists)

Access to components

in accordance with resolved type (or class)

Resolved polymorphic object

 must be type compatible with every type/class guard (constraint on guard!)

Technical problem (TP2):

 access to all extension types' definitions is needed to completely cover the inheritance tree

Type selection allows both

- run time type identification (RTTI)
- run time class identification (RTCI)

It is necessary to ensure type safety and (reasonably) good performance

- RTCI or mixed RTTI+RTCI are not expected to occur very often
- executing SELECT TYPE is an expensive operation


An RTCI scenario



"Lifting" to an extended type

- e.g., because a procedure must be executed which only works (polymorphically or otherwise) for the extended type
- remember invalid invocation of inc_datetime from earlier slide we can now write a viable version of this:

```
class(date) :: o1
: ! initialize o1
select type (o1)
class is (datetime)
  call inc_datetime(o1,.03_rk)
                     inside "class is" block:
                     • o1 is polymorphic
                      (this is what we want here!)
                     declared type of o1 is datetime
class default
  write(*,*) &
     'Cannot invoke inc_datetime on o1'
end select
```





SELECT TYPE and association



Associated alias must be used if the selector is not a named variable

 e.g., if it is a type component, or an expression

Additional restrictions:

- only one selector may appear
- the selector must be polymorphic

Example:

given the type definition

and an object o_p of that type, the RTTI for o_p%birthday is required to look like this:

```
type :: person
  class(date), allocatable :: birthday
end type
```

Polymorphism (9): A universal base class



- Denoted as "*"
 - "no declared type"
- Refers to an object that is of
 - 1. intrinsic, or
 - 2. extensible, or
 - 3. non-extensible


dynamic type

Syntax:

class(*), ... :: o_up

- an unlimited polymorphic (UP) entity
 - usual restrictions: (POINTER eor ALLOCATABLE) or a dummy argument, or both

Conceptual inheritance tree:



Polymorphism (10): UP pointer



An UP pointer can point to anything:

```
class(*), pointer :: p_up
type(datetime), target :: o_dt
real, pointer :: rval

p_up => o_dt
allocate(rval) ; rval = 3.0
p_up => rval
```

However, dereferencing ...

```
p_up => o_dt
write(*, *) p_up % yr
! will not compile
```

... is not allowed without a SELECT TYPE block (no declared type → no accessible components)

```
type(datetime), pointer :: pt

select type (p_up)
type is (datetime)
  write(*, *) p_up % yr
  pt => p_up
type is (real)
  write(*, '(f12.5)') p_up
class default
  write(*, *) 'unknown type'
end select
```

RTTI:

- can also use an intrinsic type guard in this context
- analogous for UP dummy arguments if access to data is needed

UP entities of non-extensible dynamic type





Use of this form of UP is not recommended

 Reason: different from intrinsic and extensible types, no type information is available via the object itself → SELECT TYPE always falls through to "class default"

Loss of type safety:

syntactically, it is in this case allowed to have



use this feature only if you know what you're doing
 (i.e. maintain type information separately and always check)

See examples/day2/discriminated_union for a possible usage scenario

Polymorphism (11): Allocating an UP object



Applies to

 unlimited polymorphic entities with the POINTER or ALLOCATABLE attribute

Typed allocation:

 any type may be specified, including intrinsic and nonextensible types

Sourced or molded allocation

- source or mold may be of any type (limitation to extensible type does not apply)
- the newly created object takes on the dynamic type of source or mold (same as for "regular" polymorphic objects)

Polymorphism (12): Type inquiry intrinsics



Compare dynamic types:

.TRUE. if mold is type compatible with a

- functions return a logical value
- arguments must be entities of extensible (dynamic) type, which
- can be polymorphic or non-polymorphic

Recommendation:

it may be implicitly available!

 only use if type information is not available (most typically if at least one of the arguments is UP), or if type information not relevant for the executed algorithm

That's it for today.

Following now: Exercise session 4



Object-oriented programming (II)

Binding of procedures to Types and Objects

Motivation



- Remember inc_date and inc_datetime procedures:
 - programmer decides which of the two routines is invoked
 - for an object of dynamic type date, inc_datetime cannot be invoked
- Suppose there is a desire to
 - invoke incrementation depending on the dynamic type of the object: class(date), allocatable :: o_d

```
date: o_d%increment(...) invokes inc_date
```

- datetime: o_d%increment(...) invokes inc_datetime
- This concept is also known as dynamic (single) dispatch via the object
 - cannot use F95 style generics (polymorphism forces run-time decision)

Prolegomenon: Pointers to procedures (1)





Declaration:

```
procedure(subr), pointer :: &
    pr => null()
```

- a named procedure pointer with an explicit interface ...
- ... here it is:

```
interface
  subroutine subr(x)
    real, intent(inout) :: x
  end subroutine
end interface
```

Usage:

```
real :: x
: must associate
before invocation

x = 3.0
call pr(x) ! invokes "subr"
```

Notes:

- pointing at a procedure that is defined with a generic or elemental interface is not allowed
- no TARGET attribute is required for the procedure pointed to

Pointers to procedures (2)



Functions are also allowed in this context:

```
interface
  real function fun(x)
    real, intent(in) :: x
  end function
end interface

procedure(fun), pointer :: &
    pfun => null()
```

Usage:

- this also illustrates that the target can change throughout execution (in this case to the intrinsic sin)
- some of the intrinsics get dispensation for being used like this despite being generic

Pointers to procedures (3)



Using an implicit interface 🖄

not recommended (no signature checking, many restrictions)

```
procedure(), pointer :: pi => null()
external :: targ_1, targ_2
! external, pointer :: pi => null()

procedure(), pointer :: pfi
real :: pfi, targ_2

type declaration for pfi
indicates a function pointer
```

invocations:

Procedures as type components





Two variants are supported:

object-bound procedure (OBP) and type-bound procedure (TBP)

Syntax:

"standard" type component

not obligatory

pointer to a procedure

Semantics:

 each object's %send component can be associated with any procedure with the same interface as send

```
type :: date
   :! previously defined comp.
contains
   procedure :: &
        increment => inc_date
end type
```

- Syntax:
- existing procedure
- component in contains part of type definition
- no POINTER attribute appears

Semantics:

 each object's %increment component is associated with the procedure inc_date

Restrictions on the procedure interface



... apply for both variants

First dummy argument:

This is the dummy that will usually become argument associated with the object invoking the TBP


- declared type must be same type as the type (type of the object) the procedure is bound to (the procedure pointer is a component of)
- must be polymorphic if and only if type is extensible (→ assure inheritance works with respect to any invocation)
- must be a scalar
- must not have the POINTER or ALLOCATABLE attribute

 for the type-bound case, the procedure interface has already been specified on an earlier slide


Diagrammatic representation



Object-bound procedure



Type-bound procedure (TBP)



- implementation need not be public
- increment component is
 public (even if type is opaque), unless explicitly declared private

Invocation of procedure components



Syntax is the same for the object-bound and type-bound case

 need to set up pointer association for the objectbound case before invocation

```
type(data_send_container) :: c
: ! set up desc
allocate(c%d, source = ...)
if (...) then
   c%send => my_send1
else
   c%send => my_send2
end if

call c%send(desc)
   object-bound case

assume first if branch is taken →
```

same as call my send1(c, desc)

```
type(date) :: o_d

type(datetime) :: o_dt

o_d = date(12, 'Dec', 2012)

:! also make o_dt defined

call o_d%increment(12._rk)
```

```
same as call inc_date(o_d, 12._rk)
```

Notes:

- the object is associated with the first dummy of the invoked procedure ("passed object")
- inheritance:

```
call o_dt%increment(2._rk)
```

(as things stand now) also invokes inc_date, so we haven't yet gotten what we wanted some slides earlier

Overriding a type-bound procedure



In a type extension,

an existing accessible TBP can be overridden:

with the binding above added,

```
call o_dt%increment(.03_rk)
invokes inc_datetime
```

Invoke by type component

 a class 2 name → no name space collisions between differently typed objects (with or without inheritance relation)


Invoking object: may be polymorphic or not polym.

- dynamic type is used to decide which procedure is invoked
- this procedure is unique: go up the inheritance tree until a binding is found (implicit RTCI)

Assumption:

Bold-faced types define or override TBP increment

Others don't



type may be inaccessible in invocation's scope!

Restrictions on the interface of a procedure used for overriding an existing TBP



Each must have same interface as the original TBP


- even same argument keyword names!
- if they (both!) are functions, the result characteristics must be the same
- Except the passed object dummy,
 - which must be declared class(<extended type>)
- This guarantees that inheritance works correctly together with dynamic dispatch
- In the datetime example,
 - the procedure interface of inc_datetime (see earlier slide) obeys these rules

Comment on private type-bound procedures





These cannot be overridden outside their defining module



- therefore p2 is not an overriding type-bound procedure, but a new binding that applies to all entities of class(t2)
- p2 therefore need not have the same characteristics as p

Note: compilers might get dynamic dispatch wrong in this situation, and don't handle differing interfaces (check recent releases)

Suppress overriding in extension types




- The NON_OVERRIDABLE attribute can be used in any binding
- For example, if write_date (see earlier slide) is bound to date as follows:

```
type :: date
   :! previously defined comp.
contains
  procedure :: increment => inc_date
  procedure, non_overridable :: write => write_date
end type
```

- then it is not possible to override the write TBP in any extension
- this makes sense here because it is intended that the complete inheritance tree is dealt with inside the implementation of the procedure (other rationales may exist in other scenarios)

Diagrammatic representation for overriding TBPs





Non-overridden procedures are inherited

On "SELECT TYPE" vs. "overriding TBP"



Dynamic dispatch by TBP

- TBP's should behave consistently whether handed an entity of base type or any of its extensions (Liskov substitution principle)
- example: "incrementation by (fractional) days" obeys the substitution principle
- some attention is needed to avoid violations:
 - client extends a type
 - → programmer using the interface may misinterpret intended semantics (→ documentation issue!)

- avoid bad design of extensions (analogous to side effects in functions)
- Example: derive square from rectangle (exercise)

Isolate RTTI

- to the few places where needed
 - creation of objects, I/O
- since it is all too easy to forget covering all parts of the inheritance tree
- RTCI rarely used, because TBPs fill that role
- Overriding does not lose functionality
 - parent type invocation (see left)

type(datetime) :: dtt
call dtt%date%increment(120._rk)

Array as passed object



- Passed object must be a scalar
 - therefore, arrays must usually invoke TBP or OBP elementwise
- But a type-bound procedure may be declared ELEMENTAL
 - actual argument then may be an array (remember further restrictions on interface of an ELEMENTAL procedure)
 - invocation can be done with array or array slice

```
type :: elt
   :
contains
   procedure :: p
end type
```

```
elemental subroutine p(this, x)
  class(elt), intent(inout) :: this
  real, intent(in) :: x
  :! no side effects
end subroutine
```

■ This is not feasible for the object-bound case (each elements' procedure pointer component may point to a different procedure)

Variations on the passed object: PASS and NOPASS



Pass non-first argument

- via explicit keyword specification
- example: bind procedure to more than one type

```
type :: t1
    :
contains
    procedure, &
    pass(o1) :: pf
end type

    no "=>". Why?

type :: t2
    :
contains
    procedure, &
    pass(o2) :: &
        pq => pf
end type
```

```
subroutine pf(01, x, 02, y)
  class(t1) :: o1
  class(t2) :: o2
  :
end subroutine
```

Do not pass argument at all

```
type :: t3
   :
contains
   procedure, nopass :: pf
end type
```

Invocations:

```
type(t1) :: o_t1
type(t2) :: o_t2
type(t3) :: o_t3
:
call o_t1%pf(x, o_t2, y)
call o_t2%pq(o_t1, x, y)
call o_t3%pf(o_t1, x, o_t2, y)
```

Note:

 overriding TBPs must preserve PASS / NOPASS





Properties:

- no entity of that (dynamic) type can exist
- may have zero or more components

```
type, abstract :: <type name>
   :! components, if any
[ contains
   :! type-bound procedures
]
end type
```

- declaration of a polymorphic entity of declared abstract type is permitted
- an abstract type may be an extension

Example:

```
type, abstract :: shape
end type

type, extends(shape) :: square
  real :: side
end type
```

valid and invalid usage:

Abstract Types with deferred TBPs (aka Interface Classes)





Syntax of definition

 one or more deferred bindings are added:

 cannot override a non-deferred binding with a deferred one

Deferred binding:

 described by an interface (usually abstract)

enforces that any client defining a type extension must establish an overriding binding (once you have one, it is inherited to extensions of the extension)


Extending from an interface class



```
module mod_file_handle
 use mod handle
  type, extends(handle) :: file_handle
    private
    integer :: unit
  contains
    procedure :: open => file open
  end type file handle
                                        will not compile without this override
contains
  subroutine file_open(this, info)
    class(file handle) :: this
    class(*), intent(in), optional :: info
    select type (info)
    type is (character(len=*))
      : ! open file with name info and store this%unit
      this%state = 1
    :! error handling via class default
    end select
  end subroutine
end module mod_file handle
```

Diagrammatic representation of the interface class and its realization





- Will typically use (at least) two separate modules
 - e.g., module providing abstract type often third-party-provided
- Abstract class and abstract interface indicated by italics
 - non-overridable TBP getstate() → "invariant method"

Using the interface class



Compare to "traditional" design:

- Implementation details of non-abstract type decoupled from "policy-based" design of abstract type
- Dependency inversion:
 - ideally, both clients and implementations depend on abstractions
 - in a procedural design, the type "handle" would need to contain all possible variants
 → abstraction becomes dependent on irrelevant details



Dependency Inversion with Submodules

Problems with Modules



Tendency towards monster modules for large projects

 e.g., type component privatization prevents programmer from breaking up modules where needed

Recompilation cascade effect

- changes to module procedures forces recompilation of all code that use associates that module, even if specifications and interfaces are unchanged
- workarounds are available, but somewhat clunky


Object oriented programming

- more situations with potential circular module dependencies are possible (remember TP2 on earlier slide)
- type definitions referencing each other may also occur in object-based programming





Split off implementations (module procedures) into separate files



Submodule program units



Syntax


ancestor module

Symbolic representation

```
submodule ( mymod ) smod_1
  :! specifications
contains
  :! implementations
end submodule
```

applies recursively: a descendant of smod_1 is

 sibling submodules are permitted (but avoid duplicates for accessible procedures)



Submodule specification part



Like that of a module, except

no private or public statement or attribute can appear

Reason: all entities are private

and only visible inside the submodule and its descendants

```
module mymod
  implicit none
  type :: t
   :
  end type
:
  end module
```

```
submodule ( mymod ) smod_1
  type, extends(t) :: ts
  :
  end type
  real, allocatable :: x(:,:)
:
end submodule
effectively
  private
```

Separate module procedure interface



In specification part of the ancestor module

```
module mod date
  type :: date
                                          indication that the
    :! as previously defined
                                      implementation is contained
  end type
                                           in a submodule
  interface
    module subroutine write_date (this, fname)
      class(date), intent(in) :: this
      character(len=*), intent(in) :: fname
    end subroutine
    module function create_date (year, mon, day) result(dt)
      integer, intent(in) :: year, mon, day
      type(date) :: dt
    end function
  end interface
end module
```

import statement not permitted (auto-import is done)

Separate module procedure implementation



Variant 1:

- complete interface (including argument keywords) is taken from module
- dummy argument and function result declarations are not needed

```
submodule (mod_date) date_procedures
    :! specification part
contains
    module procedure write_date
          :! implementation as shown before
    end procedure write_date
          module procedure create_date
          :! implementation as shown before
    end procedure create_date
    end procedure create_date
end submodule date_procedures
```

Separate module procedure implementation




Variant 2:

- interface is replicated in the submodule
- must be consistent with ancestor specification

```
submodule (mod date) date procedures
                                            note syntactic
  :! specification part
                                         difference to Variant 1
contains
 module subroutine write date (this, fname)
      class(date), intent(in) :: this
      character(len=*), intent(in) :: fname
    : ! implementation as shown before
  end subroutine write_date
  module function create date (year, mon, day) result(dt)
      integer, intent(in) :: year, mon, day
      type(date) :: dt
    : ! implementation as shown before
  end function create_date
end submodule date_procedures
```

Dependency inversion explained





Access to submodule entities

 can be indirectly obtained via execution of procedures declared with separate module procedure interfaces


Changes to implementations

- no dependency of program units (except descendant submodules) on these
- do not require recompilation of program units using the parent module

implementation of module procedure can access private type components due to host access to module

Exploiting dependency inversion in OO design







Generic Type-bound Procedures

Example scenario




Two existing concepts

 both support an interface of same name and function

Need to join those concepts

- which may interact in some way
- concept: multiple inheritance



■ TBP increment():

- for funds, increments amount
- for date, increments by days
- for admin_funds, both the above should work individually, and in addition it should be possible to account for the interest rate (interaction!)

These are interfaces with differing signatures!

- in principle, the funds binding will be inherited by admin_funds
- remember interface restrictions on overriding a TBP

Declaring a generic type-bound procedure



Starting point:

 the type which first declares the binding that must be generic

OCP

 may need to retrofit generic from simple TBP (easily done, at the cost of recompiling all clients)

Adding specifics to a generic in a type extension:

- three specific TBPs now can be invoked via one generic name (one inherited, two added)
- it is also allowed to bind to an inherited specific TBP

Disambiguating procedure interfaces



Implementation ...

... is inherited

... re-dispatches to
this%d%increment()

... invokes both the above, after accounting for interest

Selection of specific TBP:

- must be possible at compile time
- pre-requisite: between each pair of specifics, for at least one nonoptional argument type incompatibility is required providing two specifics which only differ in one argument, one being type compatible with the other, is not sufficient to disambiguate

Invocation of a generic TBP



```
type(admin funds) :: of
class(funds), &
      allocatable :: of_poly
allocate(admin funds :: of poly)
: ! initialize both objects
call of%increment(12, 600.)
call of%increment(17)
call of%increment(100.)
call of poly%increment(1, 2.)
```

how can this be fixed?

■ The usual TKR (type/kind/rank) matching rules apply ...

Compile-time resolution ...

```
... to inc_both()
```

... to inc_date()

... to inc_funds()

... is not possible because this interface is not defined for an entity of declared type **funds**



 a specific TBP can still be overridden i.e., compile-time resolution is only partial

See examples/day2/multiple_inheritance

Overriding a specific binding in a generic TBP



■ Further type extension (in a different module)

with a module procedure:

Invocation:

```
class(admin_funds), &
    allocatable :: o_mf

allocate(my_funds :: o_mf)
    :! initialize o_mf

call o_mf%increment(1, 23.)
```

invokes overriding procedure
 inc_my_funds because
 dynamic type is my_funds

Unnamed generic TBPs – defined operator



Example:

unary trace operator

 the NOPASS attribute is not allowed for unnamed generics

Invocation:

Rules and restrictions:

- same rules and restrictions
 (e.g., with respect to characteristics)
 as for generic interfaces and
 their module procedures
- here: procedure must be a function with an INTENT(IN) argument

Note:

 inheritance → statically typed function result may be insufficient

Unnamed generic TBPs – overloaded operator



Overloading allowed for

- existing operators
- assignment

Example:

```
type :: vector
   :! see earlier definition
contains
   procedure :: plus1
   procedure :: plus2
   procedure, pass(v2) :: plus3
   generic, public :: &
        operator(+) => plus1, plus2, plus3
end type matrix
```

Specifics:

```
function plus1(v1, v2)
  class(vector), intent(in) :: v1
  type(vector), &
             intent(in) :: v2
  type(vector) :: plus1
  : ! implementation omitted
end function
function plus2(v1, r)
  class(vector), intent(in) :: v1
  real, intent(in) :: r(:)
  type(vector) :: plus2
  : ! implementation omitted
end function
function plus3(r, v2)
class(vector), intent(in) :: v2
  real, intent(in) :: r(:)
  type(vector) :: plus3
  : ! implementation omitted
end function
```

Using the overloaded operator



```
type(vector) :: w1, w2
real :: r(3)

w1 = vector( [ 2.0, 3.0, 4.0 ] )
w2 = vector( [ 1.0, 1.0, 1.0 ] )
r = [ -1.0, -1.0, -1.0 ]

invokes plus1( (w1), (w2) )

w2 = w1 + w2
w2 = w2 + r
w2 = r + w1

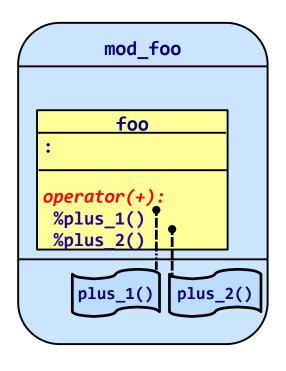
invokes plus2( (w2), (r) )
```

Remaining problem:

- how to deal with polymorphism –
- for an extension of vector, the result usually should also be of the extended type
- but: function result must be declared consistently for an override

Diagrammatic representation of generic TBPs





following now: Exercise session 6

Use italics to indicate generic-ness

- provide list of specific TBPs as usual
- overriding in subclasses can then be indicated as previously shown



Nonadvancing I/O

Reminder on error handling for I/O



- ☐ An I/O statement may fail: Examples:
 - opening a non-existing file with status='OLD'
 - reading beyond the end of a file
- Without additional measures:RTL will terminate the program
- Prevent termination via: user-defined error handling
 - specify an iostat and possibly iomsg argument in the I/O statement
 - use of err / end / eor = <label>
 is also possible but is legacy!
 - → do not use in new code!!

■ iostat=ios specification

ios (scalar default integer) will be:

- negative if end of file detected,

- positive if an error occurs,

- zero otherwise

☐ iomsg=errstr specification

errstr (default character string of sufficient length) supplied with appropriate description of the error if iostat is none-zero

☐ Use intrinsic logical functions:

```
is_iostat_end(ios)
is_iostat_eor(ios)
```

to check iostat-value of I/O operation

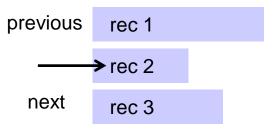
for EOF (end of file) or EOR (end of record)

condition

Nonadvancing I/O (1)



Allow file position to vary inside a record:

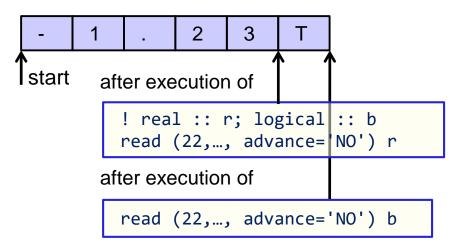


Syntactic support:

 ADVANCE specifier in formatted READ or WRITE statement

on record No. 2

read with '(f5.2)', '(11)' - each square is 1 character (byte)



if a further READ statement is executed, it would abort with an end-of-record condition.

retrieve iostat-value (default integer) via iostat specifier: allows handling by user code and positions connection at beginning of next record:

```
read (...,advance='NO',iostat=ios) ...
if(is_iostat_eor(ios)) ...
```

Nonadvancing I/O (2)



□ Reading character variables

 the SIZE specifier allows to determine the number of characters actually read

```
character(len=6) :: c
integer :: sz
:
!read chars from file into string:
read(23,fmt='(a6)',advance='NO',&
    pad='YES', iostat=ios, size=sz) c
! Set remaining chars to
! a non-blank char if EOR occurs:
if (is_iostat_eor(ios)) c(sz+1:)='X'
```

 mainly useful in conjunction with EOR (end-of-record) situations

■ Nonadvancing writes

 usually used in form of a sequence of nonadvancing writes, followed by an advancing one to complete a record

☐ Final remarks

- nonadvancing I/O may not be used in conjunction with namelist, internal or list-directed I/O
- several records may be processed by a single I/O statement also in non-advanced mode
- format reversion takes precedence over non-advancing I/O



Object-oriented I/O Facilities: User defined derived type I/O

I/O for derived data types



Non-trivial derived data type

```
type :: list
  character(len=:), &
        allocatable :: name
  integer :: age
  type(list), pointer :: next
end type
```

- Perform I/O using suitable module procedures
- Disadvantages:
 - recursive I/O disallowed
 - I/O transfer not easily integrable into an I/O stream
 - defined by edit descriptor for intrinsic types and arrays
 - or sequence of binary I/O statement

F03

 enables binding a subroutine to an I/O list item of derived type

```
type(list) :: o_list
: ! set up o_list
write(unit, fmt='(dt ...)', ...) &
    o_list
```

- example shows formatted output
- bound subroutine called automati-cally when edit descriptor DT is encountered
- other variants are enabled by using generic TBPs or generic interfaces
- can use recursion for hierarchical types

Binding I/O subroutines to derived types



- Interface of subroutines is fixed
 - with exception of the passed object dummy
- Define as special generic type bound procedure

```
type :: foo
   :
contains
   :
   generic :: read(formatted) => rf1, rf2
   generic :: read(unformatted) => ru1, ru2
   generic :: write(formatted) => wf1, wf2
   generic :: write(unformatted) => wu1, wu2
end type
```

- genericness refers to rank, kind parameters of passed object
- Define via interface block

```
interface read(formatted)
  module procedure rf1, rf2
end interface
```

DTIO module procedure interface





```
subroutine rf1(dtv,unit,iotype,v_list,iostat,iomsg)
subroutine wu1(dtv,unit, iostat,iomsg)
```

□ dtv:

- scalar of derived type
- may be polymorphic
- of suitable intent

☐ unit:

- integer, intent(in) describes I/O unit or negative for internal I/O
- iotype (formatted only):
- character, intent(in) 'LISTDIRECTED',
 'NAMELIST' or 'DT'//string
 see dt edit descriptor

v_list (formatted only):

 integer, intent(in) - assumed shape array see dt edit descriptor

☐ iostat:

- integer, intent(out) scalar, describes error condition
- iostat_end / iostat_eor / zero if all OK

☐ iomsg:

• **character(*)** - explanation for failure if iostat nonzero

Limitations for DTIO subroutines



I/O transfers to other units than unit are disallowed

- I/O direction also fixed
- Exception:internal I/O is OK (and commonly needed)

Use of the statements

- open, close, rewind
- backspace, endfile

is disallowed

File positioning:

- entry is left tab limit
- no record termination on return
- positioning with
 - rec=... (direct access) or
 - pos=... (stream access)


is disallowed

Writing formatted output:

DT edit descriptor



Example:



- Both iotype and v_list are available to the programmer of the I/O subroutine
 - determine further parameters of I/O as programmer sees fit
- Note:
 - inside a formatted DTIO procedure ("child I/O"), I/O is nonadvancing (no matter what you specify for ADVANCE)

Example: Formatted DTIO on a linked list



■ Here: type definitions and DTIO-procedure implementations inside same module, e.g.:

```
module mod_list

type :: list
   integer :: age
   character(20) :: name
   type(list), pointer :: next

contains

generic :: &
   write(formatted) => wl

end type list

contains

: ! to be continued
```

```
: ! module mod list continued
  recursive subroutine wl &
    (this,unit,iotype,vlist,iostat,iomsg)
    class(list), intent(in) :: this
    integer , intent(in) :: unit, vlist(:)
    character(*), intent(in) :: iotype
    integer, intent(out):: iostat
    character(*)
                       :: iomsg
    ! .. Locals
    character(len=12) :: pfmt
    if (iotype /= 'DTList') return
    if (size(vlist) < 2) return</pre>
    ! internal IO to generate format descriptor
    write(pfmt, '(a,i0,a,i0)') &
                '(i',vlist(1),',a',vlist(2),')'
    write(unit, fmt=pfmt, iostat=iostat) &
                this%age, this%name
    if (iostat /= 0 ) return
    if (associated(this%next)) call wl &
    (this%next,unit,iotype,vlist,iostat,iomsg)
  end subroutine
    !:other implementations ...
end module
```

Example (cont'd): Client use



Client use formatted DTIO

```
type(list), pointer :: mylist
  ! : set up mylist
  ! : open formatted file to unit

write(unit,fmt='(dt "List" (4,20) )', &
        iostat=is) mylist
! : close unit and destroy list
```

Client use unformatted DTIO

```
type(mydt) :: o_mydt
! : unformatted writing (also) bound to mydt
! : open unformatted file to unit 21
write(21[, rec=...]) o_mydt
```

Final remarks: Unformatted DTIO

- bound subroutine with shorter argument list
- is automatically invoked upon execution of write statement
- additional arguments (e.g. record number) only specifiable in parent data transfer statement

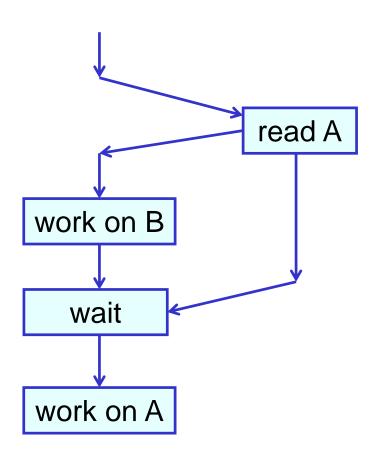


Asynchronous I/O

Basic concept



Flow diagram



Implementation

```
real, dimension(100000) :: a, b
open(20,...,asynchronous='yes')
...
read(20,asynchronous='yes') a
! do work on something else
:
wait(20)
! do work with a
... = a(i)
```

Ordering requirements

- apply for a sequence of data transfer statements on the same I/O unit
- but not for data transfers to different units

Conditions for asynchronous execution



Necessary conditions

OPEN statement

```
open(unit=<number>, ...,asynchronous='yes')
```

- prerequisite for performing asynchronous I/O statements on that unit
- READ or WRITE statements

```
[read|write](unit=<number>,..., &
    asynchronous='yes',id=<integer_tag>) <data_entity>
```

 ID specifier allows to assign each individual statement a tag for subsequent use

Actual asynchronous execution

- is at processors discretion
- is most advantageous for large, unformatted transfers

The WAIT statement



Block until asynchronously started I/O statement has completed

```
wait(unit=<number> [, &
    id=<integer_tag>, &
    end=<label>, &
    eor=<label>, &
    err=<label>, &
    iostat=<integer_status>])
```

- if ID tag present → wait only for tagged statement
- else wait for all outstanding I/O statements on that unit

Implication:

 can have multiple outstanding asynchronous I/O statements to the same unit

Implicit WAIT is incurred:

- by BACKSPACE
- by REWIND
- by ENDFILE
- possibly by INQUIRE
- by CLOSE

Orphaned WAIT

- with respect to unit or id
- has no effect

Non-blocking execution



Option for check of I/O completion

extension of INQUIRE statement

```
inquire(unit=<number>, pending=<logical>, id=<integer_tag>)
```

- PENDING specifier returns .true.
 if operation tagged by ID is not yet complete
- if no ID present, all outstanding I/O statements must be complete
- PENDING specifier returns .false.
 if operation tagged by ID is complete
- refers to completion status of all outstanding I/O statements if no ID present
- a return value .false. implies a WAIT (i.e. an implementation may decide to wait for completion while the INQUIRE executes)

Affector entities



- Entity in a scoping unit
 - item in an I/O list
 - item in a NAMELIST
 - SIZE= specifier
- associated with an asynchronous I/O statement
- Constraints on affectors:
 - must not be redefined,
 - become undefined, or
 - have pointer association status changed
- while I/O operation on it is pending

While asynchronous input is pending

 affector must not be referenced or associated with a VALUE dummy argument

Affectors and Optimization (1)



Recall prototypical case:

```
open(20,...,asynchronous='yes')
...
read(20,asynchronous='yes') a
:
! compiler may not prefetch "a" here
wait(20)
... = a(i)
```

- asynchronous I/O puts constraints on code movement by the compiler
- all affectors automatically acquire the ASYNCHRONOUS attribute in the above case
 - once acquired in a scoping unit, will propagate
 - all subobjects of an affector also have the attribute

Affectors and Optimization (2)



However consider

```
subroutine read async(unit,id,this)
  integer, intent(in) :: unit
  integer, intent(out) :: id
 type(...), intent(out) :: this
  read(unit=iu,id=id, &
       asynchronous='yes') this
end subroutine
subroutine work(unit,id,this)
  integer, intent(in) :: unit
  integer, intent(in) :: id
 type(...), intent(in), &
             asynchronous :: this
 wait(unit, id=id)
 ... = this
end subroutine
```

 need explicit attribute to suppress code motion

Consider further the call sequence

```
type(...), asynchronous :: this
: ! Open unit
call read_async(unit,id,this)
:
call work(unit,id,this)
... = this
```

 due to intent(in) in work() compiler could move loads of this before call to work()

again, need explicit **ASYNCHRONOUS** attribute



Performance considerations for using I/O

Expected types of I/O



1. Configuration data

- usually small, formatted files
- parameters and/or meta-data for large scale computations

2. Scratch data

- very large files containing complete state information
- required e.g., for checkpointing/restarting
- → rewrite in regular intervals
- throw away after calculation complete

3. Data for permanent storage

- result data set
- for post-processing
- to be kept (semi-) permanently
- archive to tape if necessary
- may be large, but not (necessarily) complete state information

Which file system(s) should I use?



For I/O of type 1:

- any will do
- if working on a shared (possible parallel) file system:

Beware transaction rates

- → OPEN and CLOSE stmts may take a long time
- → do not stripe files

For I/O of type 2 or 3:

- need a high bandwidth file system
- → parallel file system with block striping
- large file support nowadays standard

What bandwidths are available?

- normal SCSI disks~100 MByte/s
- NAS storage arrays at LRZ: up to 2 GByte/s
- SuperMUC storage arrays:
 up to 150 GByte/s



- aggregate for all nodes
- single node can do up to 2 GB/s (large files striped across disks)
- → writing the memory content of system to disk takes ~40 minutes

I/O formatting issues various ways of reading and writing



better performance

Formatted I/O

list directed

with format string

can be static or dynamic

Unformatted I/O

sequential

direct access


can also be formatted

I/O access patterns

better

performance

- by implicit loop
- by array section
- by complete array



write(...) a

I/O performance for implicit DO loops



Improve performance by

- imposing correct loop order (fast loop inside!)
- more important: writing large block sizes

Large blocks, but wrong order.
On some platforms this may give a performance hit

→ re-copy array or reorganize data

- proper tuning
 - > performance may exceed that for array sections

Discussion of unformatted I/O properties



- No conversion needed
 - saves CPU time
- No information loss
- Needs less space on disk
- File not human-readable
 - binary
 - Fortran record control words
 - possible interoperability problems with I/O in C
 - convert to Stream I/O
- Format not standardized
 - in practice much the same format is used anyway
 - exception big/little endian issues
 - solvable if all data types have same size

Support for little/big endian conversion by Intel compiler

- enable at run time
- suitable setting of environment variable F_UFMTENDIAN
- example:

```
export F_UFMTENDIAN="little; big:22"
```

will set unit 22 **only** to big-endian mode (little endian is default)

- performance impact??
- other compilers might need:
 - changes to source or
 - compile time switch

I/O and program design



Except for debugging or informational printout

- try to encapsulate I/O as far as possible
 - → each module has (as far as necessary) I/O routines related to it's global data structures
 - → mapping of file names should reflect this
- write extensibly, i.e.: use a generic interface which can then be applied to an extended type definition
 - in fact module internal code can usually be re-used
 - keep in mind: performance issues may crop up if code used outside its original design point

Additional documentation requirement

description of structure of data sets needed



IEEE Arithmetic and IEEE Floating Point Exception Handling

IEEE-754



ISO-IEC standard for binary floating point processing

Defines

- floating point representations
- prescriptions for conforming +,-,*,/,sqrt (portable FP programming)
- rounding modes
- exceptions and exception handling mechanisms

Standard CPUs have hardware support for (most of) the above

- Reference: David Goldberg's article
 - What Every Computer Scientist Should Know about Floating-Point Arithmetic
- Further information available at
 - http://grouper.ieee.org/groups/754/

Intrinsic modules (1)



IEEE support in Fortran

• three intrinsic modules → subset support


IEEE features

contains a number of constants of type ieee_feature_type

```
use, intrinsic :: &
  ieee_features, only : ieee_divide
```

support IEEE conforming divide

- effect as for a compiler switch
- applies for complete scoping unit
- possible performance impact → use an only clause to limit effect
- if a feature unsupported → compilation fails



List of constants in ieee_features



Default settings:

may be an arbitrary subset

Named constant	Effect of access in scoping unit for at least one kind of real	
ieee_datatype	must provide IEEE arithmetic	
ieee_denormal	must support denormalized numbers	
ieee_divide	must support IEEE divide	
ieee_halting	must support control of halting	
ieee_inexact_flag	must support inexact exception	
ieee_inf	must support -∞ and +∞	
ieee_invalid_flag	must support invalid exception	
ieee_nan	must support NaN	
ieee_rounding	must supportcontrol of all four rounding modes	
ieee_sqrt	must support IEEE square root	
ieee_underflow_flag	must support underflow exception	

Intrinsic modules (2)




IEEE exceptions

- definitions of types, flags and procedures for exception handling
- note: no support for (user-defined) handler callback functions

IEEE arithmetic

- definitions of classes of floating point types
- definitions of rounding modes; functions for setting and getting these modes
- Using ieee_arithmetic implies
 use of ieee_exceptions



Models for integer and real data



Numeric models for integer and real data

 $i = s \times \sum_{k=1}^{q} w_k \times r^{k-1}$

integer kind is defined by

- positive integer q (digits)
- integer r > 1 (normally 2)

integer value is defined by

- sign $s \in \{\pm 1\}$
- sequence of $w_k \in \{0, ..., r-1\}$

base 2 → "Bit Pattern"

integers are not dealt with through the IEEE facilities

fractional part

$$x = b^e \times s \times \sum_{k=1}^p f_k \times b^{-k} \quad \text{or } \mathbf{x} = \mathbf{0}$$

real kind is defined by

- positive integers p (digits),
 b > 1 (base, normally b = 2)
- integers e_{min} < e_{max}

real value is defined by

- sign $s \in \{\pm 1\}$
- integer exponent $e_{min} \le e \le e_{max}$
- sequence of $f_k \in \{0, ..., b-1\}$, f_1 nonzero

Inquiry intrinsics for model parameters



digits(x)	for real oder integer x, returns the number of digits (p, q respectively) as a default integer value.	<pre>minexponent(x), maxexponent(x)</pre>	for real x, returns the default integer \mathbf{e}_{min} , \mathbf{e}_{max} respectively
precision(x)	for real or complex x, returns the default integer indicating the decimal precision (=decimal digits) for numbers with the kind of x.	radix(x)	for real or integer x, returns the default integer that is the base (b, r respectively) for the model x belongs to.
range(x)	for integer, real or complex x, returns the default integer indicating the decimal exponent range of the model x belongs to.		


Inquiry intrinsics for model numbers



Example representation: $e \in \{-2, -1, 0, 1, 2\}, p=3$

purely illustrative!

• look at first positive numbers (spacings $\frac{1}{32}$, $\frac{1}{16}$, $\frac{1}{8}$ etc.)



rrspacing(x) = abs(x) / spacing(x)

largest representable number: ⁷/₂
 (beyond that: overflow)
 huge(x)

Mapping fl: $\mathbb{R} \ni x \to fl(x)$

- to nearest model number
- maximum relative error

$$fl(x) = x \cdot (1+d), |d| < u$$

... more realistic models



Typically used representations: IEEE-754 conforming

matched to hardware capabilities

real kind	dec. digits	base 2 digits	dec. exponent range	base 2 exponent range
default	6	24	10 ⁻³⁷ 10 ⁺³⁸	-125 +128
extended	15	53	10 ⁻³⁰⁷ 10 ⁺³⁰⁸	-1021 + 1024

Negative zero:

- hardware may distinguish from positive zero
- e.g., rounding of negative result toward zero retains sign,
- e.g., I/O operations (sign stored in file)

Closure issues (1): Rounding



Arithmetic operations:

- result typically outside the model
- requires rounding
- implementation dependency, but all good ones adhere to "standard model"

$$fl_{op}(x,y) = (x \ op \ y) \cdot (1+d),$$

 $| \ d \ | \le u; op = + / -,*,/.$

precision achieved in IEEE algorithms by using guard digits

There exist relevant algorithms for which less strict models cause **failure!**

Rounding modes:

- modify exact result to become a representable number
- nearest: to nearest representable value (NRV)
- to-zero: go toward zero to NRV
- up: go toward +∞ to NRV
- down: go toward -∞ to NRV

Note:

- division a/b executed as a*(1/b) may not be IEEE conforming (roundoff)
- conversely: enforcing IEEE conformance may have performance impact

Closure issues (2): Rounding



- IEEE-754 rounding modes
 - all fulfill the model from the previous slide
- Named constants in module ieee_arithmetic

```
ieee_nearest
ieee_to_zero
ieee_up
ieee_down
ieee_other
all of type ieee_round_type
```

Ask for full rounding support:

```
use, intrinsic :: ieee_features, only : ieee_rounding
```

Example program illustrating rounding



```
use, intrinsic :: ieee arithmetic, only : ieee_nearest, ieee_up, &
     & ieee down, ieee set rounding mode, ieee support rounding
real :: a, d
                                                replace XXX by desired mode
integer :: i
d = 0.232
if (ieee_support_rounding(ieee_XXX,a)) then
  write(*,fmt='(''Round XXX:'')')
  call ieee set rounding mode(ieee XXX)
  a = 1.5
  do i = 0, 4
      a = a / d
     write(*, fmt='(f13.6)') a
   end do
else
  write(*, fmt='(''Rounding mode ieee XXX unsupported'')')
end if
```

Produces output:

Pound noanoct

Round up:
6.465518
27.868612
120.123337
517.773071
2231.780 <mark>762</mark>

Round down:
6.465517
27.868608
120.123314
517.772888
2231.779 <mark>541</mark>

Control propagation of rounding error



Strategy 1:

- choose most appropriate rounding
- randomized may be best, but is expensive!

Strategy 2: Interval arithmetic

[a,b] (ordered interval in \mathbf{R}) $\mu([a,b]) := b - a$

define operations

$$[a,b] + [c,d] := [a+c,b+d]$$

 $[a,b] * [c,d] := [ac,bd] if a,c>0 etc.$

- keep actual value within small interval
- "small" may become difficult:

$$\mu([a,b]^*[c,d]) = d^*\mu([a,b]) + a^*\mu([c,d])$$

Implementation of IA

 may want to use rounding to guarantee enclosure

Note:

- there exist multiple variants of IA
- a standardization effort (outside Fortran) is under way

Interval software

INTLIB library /
 interval_arithmetic Fortran
 module from R. Baker Kearfott's
 page at
 http://interval.louisiana.edu/

Obtaining the current rounding mode



Intrinsic module procedure:

- entities of opaque typeieee round type
- the only allowed operations are assignment, ==, /=

Remember:

- rounding error and truncation error are two different things
- the latter usually arises from finite approximation of a representation of a function; explicit truncation error terms can sometimes (but not always) be established

Closure issues (3): special FP numbers




Three variants:

- ∞, -∞
- NaN (signaling or quiet)
- denormal numbers (f₁ = 0 and e=e_{min})

From which operations?

- 1.0 / 0.0 (or -1.0 / 0.0), more general: numbers exceeding HUGE(x) (or smaller than -HUGE(x))
- 0.0 / 0.0, SQRT(-1.0), more general: invalid operations
- gradual underflow



Production of special FP numbers triggers exceptions

IEEE arithmetic: Inquiry functions for real and complex types



logical function specification	Description	
<pre>ieee_support_datatype([x])</pre>	supports at least a subset of IEEE arithmetic (operations, rounding mode, kind, some intrinsics)	
<pre>ieee_support_denormal([x])</pre>	supports IEEE denormalized numbers	
<pre>ieee_support_divide([x])</pre>	supports IEEE conformant division	
<pre>ieee_support_inf([x])</pre>	supports IEEE infinity facility	
<pre>ieee_support_nan([x])</pre>	supports IEEE Not-a-Number facility	
<pre>ieee_support_rounding(rd_value [,x])</pre>	supports specified rounding mode as well as setting via intrinsic. Additional argument is type(ieee_round_type), intent(in) :: rd_value	
<pre>ieee_support_sqrt([x])</pre>	supports IEEE conformant square root	
<pre>ieee_support_standard([x])</pre>	supports all IEEE facilities within ieee_arithmetic	
<pre>ieee_support_underflow_control([x])</pre>	supports control of IEEE underflow mode via intrinsic	

all functions return a result of type logical and take (at least) an optional argument of type real

IEEE arithmetic:



Underflow handling & real kinds

```
subroutine ieee_get_underflow_mode(gradual)
```

- logical, intent(out) :: gradual
- true. is returned of gradual underflow is in effect

```
subroutine ieee_set_underflow_mode(gradual)
```

- logical, intent(in) :: gradual
- true. lets gradual underflow come into effect

```
integer function ieee_selected_real_kind([p] [,r])
```

- same functionality as selected_real_kind()
- but returns a kind value for which ieee_support_datatype(x) is true

The five exceptions defined in IEEE-754



Treated via (hardware) flags

 these are sticky → signal is maintained until explicitly reset by client

- 1. overflow: exact result of operation too large for model
- 2. divide_by_zero
- **3. invalid**: operations like ∞*0, 0/0, whose result is a signaling NaN

- **4. underflow**: result is finite, but too small to represent with full precision within model
- → store best result available
- **5. inexact**: exact result cannot be represented without rounding
- → will be very common


Exception handling

- check exception flags and write code to deal appropriately with the situation
- halting execution immediately or delayed (arbitrarily)

Fortran facilities – mapping the exception flags



Constants of opaque type ieee_flag_type



- Each flag can essentially have two states
 - signaling or non-signaling

Example: Division by zero



Without using IEEE facilities

```
subroutine invert(z)
  real, intent(inout) :: z
  z = 1.0 / z
end subroutine exception may be
  triggered here
```

exception may cause halting

 → control behaviour via
 compiler switch if possible

 STOP statement: if it is missing, it is implementationdependent whether any exception status is reported at termination

Example cont'd: Division by zero



With use of IEEE facilities

```
subroutine invert(z)
 use, intrinsic :: &
        ieee exceptions, only: &
        ieee divide by zero, &
        ieee set flag, &
        ieee get flag
 real, intent(inout) :: z
                              clear exception
 logical :: sig
                                 on entry
 call ieee_set_flag( &
   ieee divide by zero, .false.)
 z = 1.0 / z
 call ieee_get_flag( &
   ieee divide_by_zero, sig)
 if (sig) write(*,*) &
       'FPE div by zero signaled'
 call ieee set flag( &
   ieee divide by zero, .false.)
end subroutine
```

Invoking program

```
use, intrinsic :: ieee features, &
     only: ieee halting
use, intrinsic :: ieee_arithm@ssure halting
                            mode can be set
call ieee set halting mode( &
      ieee_divide_by_zero, .false.)
if (.not. ieee support flag( &
    ieee divide by zero, xz)) then
  stop 'FPE div by zero unsupp.'
end if
xz = 0.0
write(*,fmt='(''Invert xz='', &
        E10.3)') xz
call invert(xz)
write(*,fmt='(''Result is :'', &
        E10.3)') xz
```

Output:

```
Invert xz= 0.000E+00
FPE div by zero signaled
Result is : Infinity
```

clear exception

on exit

Handling of exceptions (1): Halting mode



```
logical function ieee_support_halting(flag)
```

- type(ieee_flag_type), intent(in) :: flag
- returns .true. if specified flag supports halting

```
subroutine ieee_get_halting_mode(flag, halting)
```

- type(ieee_flag_type), intent(in) :: flag
- logical, intent(out) :: halting
- halting is set to .true. if flag signaling causes halting

```
subroutine ieee_set_halting_mode(flag, halting)
```

- type(ieee_flag_type) :: flag
- logical, intent(in) :: halting

- can be arrays
- if halting is .true., the flag signaling will cause halting
- may only be called if ieee_support_halting(flag) is .true.

Handling of exceptions (2): Flag support



```
logical function ieee_support_flag(flag [,x])
```

- type(ieee_flag_type), intent(in) :: flag
- returns .true. if specified flag supported [for kind of x]

```
subroutine ieee_get_flag(flag, flag_value)
```

- type(ieee_flag_type), intent(in) :: flag
- logical, intent(out) :: flag_value
- flag_value is set to .true. if specified flag signals

```
subroutine ieee_set_flag(flag, flag_value)
```

- type(ieee_flag_type) :: flag
- logical, intent(in) :: flag_value


- can be arrays
- if flag_value is .false., specified flag will be set to quiet
- will only set signaling if ieee_support_flag(flag) is .true.

Disambiguate two exceptions of the same kind



Background:

- manipulating flags is an expensive operation
- trace exception for complete code blocks with many FP operations, and only a few potential failure points
- more than one failure point with the same exception
 - → need to record exceptions for disambiguation




Handling of exceptions (3):

Saving and restoring FP state



Opaque derived type

- ieee_status_type
- Object of that type stores
 - all exception flags
 - rounding mode
 - halting mode
 - complete FP state
- It cannot be directly accessed for information



unrecoverable

error → abort

Two transfer routines

- type(ieee_status_type), &
 intent(out) :: status_value
- reads FP state into status_value

- type(ieee_status_type), &
 intent(in) :: status_value
- write FP state back to flags/registers
- then, use ieee_get_flag() etc to retrieve information

execute slower but safe algorithm exec. sequence

y status(2)

y status(1)

Determine IEEE class



Named constants of type

ieee_class_type

- ieee_signaling_nan
- ieee quiet nan
- ieee_negative_inf
- ieee_negative_normal
- ieee negative denormal
- ieee_negative_zero
- ieee_positive_zero
- ieee_positive_denormal
- ieee_positive_normal
- ieee_positive_inf
- ieee_other_valueunidentifiable (e.g., via I/O)

This opaque type supports

- assignment
- Elemental intrinsics:

 return the class a real number belongs to.

```
logical function
ieee_is_ finite
nan
negative
normal
(x)
```

- identify FP class;
- use e.g., to disambiguate multiple exceptions.

Further IEEE class functions



```
logical function &
    ieee_unordered(x,y)
```

 returns .true. if one or both arguments are NaN

- produce special number values (Infinity, NaN, denormals) if supported
- invoking the intrinsic should not trigger a FP signal
- values are processor-dependent, but are the same between different invocations with the same argument

Further IEEE intrinsics



in part mirror-images of standard intrinsics

```
nearest() → ieee_next_after()
```

many elemental (i.e., applicable to arrays as well as scalars)

 please consult standard (section 14.11) for specification of these intrinsics

Example programs:

see examples/day4/ieee

This is it for today! tomorrow morning: Exercise session 8



Parameterized Derived Types: Introduction



- So far we have seen three important concepts related to OOP-paradigm: inheritance, polymorphism and data encapsulation
- ☐ Here we add another concept:
 - Concept of a parameterized derived type
- We know the concept already, have a look at object declarations of intrinsic type:

```
    All intrinsic types are actually
parameterized with the kind
parameter (intrinsic types: integer,
real, complex, logical, character)
```

- Objects of type character are additionally parameterized with the len parameter
- We extend the concept to derived types, e.g.:

```
! scalar of type real
! with non-default kind:
real(kind=real32) :: a
! array of integer numbers
! with non-default kind parameter
integer(kind=int64) :: numbers(n)
!character of default kind
!with deferred length parameter:
character(len=:), allocatable :: path
```

```
!define parameterized type:
type pmatrixT(k,r,c)
  integer, kind :: k
  integer, len :: r,c
  real(kind=k) :: m(r,c)
end type
!declare an object of that type
type(pmatrixT(real64,30,20)) :: B
```

Kind and Length Parameters



☐ F2003 permits type parameters of derived type objects.

Two varieties of type parameters exist:

- kind parameters, must be known at compile time
- Length parameter which are also allowed to be known only during runtime

```
!kind parameters from intrinsic module
use iso_fortran_env, only: real32, real64
!define parameterized type:
type pmatrixT(k,r,c)
   integer, kind :: k
   integer, len :: r,c
   real(k) :: m(r,c)
end type
!: declare an object of that type
type(pmatrixT(real32,30,20)) :: A
type(pmatrixT(real64,10,15)) :: B
```

- Type parameters are declared the same way as usual DT-components with the addition of specifying either the kind or len attribute
 - k here resolves to compiletime constant real32 (for A) and real64 (for B)
 - r,c could be deferred but here resolves to literal constants 30,20 (A) and 10,15 (B)



Parameterized Derived Type vs. Conventional Derived Type

```
module mod pmatrix
!define parameterized type:
  type pmatrixT(k,r,c)
    integer, kind :: k
    integer, len :: r,c
    real(k) :: m(r,c)
  end type
contains
subroutine workona pmat32(cs,rs)
    integer :: cs,rs
    type(pmatrixT(real32,cs,rs)) :: M
    !M%m(:,:) = ...
  end subroutine
subroutine workona pmat64(cs,rs)
    integer :: cs,rs
    type(pmatrixT(real64,cs,rs)) :: M
    !M%m(:,:) = ...
  end subroutine end module
end module
               advantage:
               1 single type definition
               2 dynamic data in component
               without allocatable or pointer
! client use attribute
call workona_pmat32(20,30)
call workona pmat64(20,30)
```

```
module mod matrix
  type matrix32T
    real(real32),allocatable:: m(:,:)
  end type
 type matrix64T
    real(real64),allocatable:: m(:,:)
  end type
contains
  subroutine workona mat32(cs, rs)
    type(matrix32T) :: M
    allocate(M%m(cs,rs))
    !M%m(:,:) = ...
  end subroutine
 subroutine workona mat64(cs, rs)
    type(matrix64T) :: M
    allocate(M%m(cs,rs))
    !M\%m(:,:) = ...
                    disadvantage:
  end subroutine
end module
                    1 two type definitions
                    2 dynamic data only
                    through allocatable or
                    pointer attribute
! client use
call workona mat32(20,30)
call workona_mat64(20,30)
```

Inquire Type parameters



■ Type parameters of a parameterized object can be accessed directly using the component selector

```
!type definition as in previous example
type(pmatrixT(real64,cols,rows)) :: A

write(*,*) A%k
write(*,*) A%c
write(*,*) A%r
do i = 1,A%c
    do j = 1,A%r
        A%m(i,j) = ...
enddo
enddo
```

☐ However, type parameters cannot be directly modified, e.g.:

```
type(pmatrixT(real64,cols,rows)) :: A
A%k=real32 ! invalid
A%c=8 ! invalid
A%r=12 ! invalid
```

Assumed Type Parameters



■ Let's pass a parameterized object into a subroutine

```
!type definition as in previous example
type(pmatrixT(real64,20,30)) :: A
type(pmatrixT(real64,10,20)) :: B

call proc_pmat(A)
call proc_pmat(B)
```

- ☐ The len parameter can be assumed from the actual argument using the *-notation
- NOTE! The kind parameter cannot be assumed!
 - But dealing with the (few) different kind parameters of interest is potentially more manageable than having to additionally deal with all len-parameter combinations
- NOTE! Type parameters cannot be assumed if dummy object has the allocatable or pointer attribute

```
module mod pmatrix
 !: definitions as before
 interface proc pmat
  module procedure :: proc pmat32, &
            proc pmat64
end interface
contains
  subroutine proc pmat64(M)
    ! dummy with assumed len parameters:
    type(pmatrixT(real64,*,*)) :: M
    do i = 1,M%c
       do j = 1,M%r
          M\%m(i,j) = ...
       enddo
     enddo
  end subroutine
 subroutine proc pmat32(M)
   type(pmatrixT(real32,*,*)) :: M
end subroutine
 subroutine otherwork_pmat64(M1,M2)
   type(pmatrixT(real64,*,*)), &
                 allocatable :: M1 !invalid
   type(pmatrixT(real64,*,*)), &
                     pointer :: M2 !invalid
 end subroutine
end module
```

Deferred Type Parameters



■ Using the colon notation we may declare objects of parmeterized derived type with deferred len-parameter if they have the pointer or allocatable attribute

```
!type definition as in previous example
type(pmatrixT(real32,:,:)), allocatable :: A, B
type(pmatrixT(real32,:,:)), pointer :: P
type(pmatrixT(real32,5,8)) :: M_5_8

allocate(type(pmatrixT(real32,15,10)::A)
P => M_5_8
allocate(B, source=P) !B allocated B%r=5, B%c=8
```

☐ The previous invalid code (assumed len parameter for allocatable dummy object) can be corrected using deferred len parameters using colonnotation for passed dummy objects with allocatable or pointer attribute

```
module mod_pmatrix
!: definitions as before
contains
!:
   subroutine otherwork_pmat64(M1,M2)
      type(pmatrixT(real64,:,:)), allocatable :: M1 ! valid
      type(pmatrixT(real64,:,:)), pointer :: M2 ! valid
   end subroutine
!:
end module
```

Default Type Parameters



☐ It is possible to define default parameters for a parameterized derived type

- ☐ You may specify only a subset of parameters and/or out of order but it requires to use keyword notation to correctly associate each actual parameter with the right type-parameter
- This also applies to deferred or assumed len declarations:

```
type(pmatrixT(k=real32,c=*,r=*)) :: M_assumed
type(pmatrixT(c=:,r=:,k=real32)), allocatable :: M_deferred
type(pmatrixT(c=:,r=:,k=real32)), pointer :: M_pointer
```

Inheritance and polymorphism



- ☐ It is possible to inherit properties from an existing base type via type extension
- Extended types may add additional kind and/or len parameters for subsequent component declarations

```
type mat_aT(k,r,c)
   integer, kind :: k=real64
   integer, len :: r=1,c=1
end type
type,extends(mat_aT) :: mat_rT
   real(k) :: m(r,c)
end type
type,extends(mat_aT) :: mat_crT(k2,m1)
   real(k) :: m(r,c)
   integer, kind :: k2=int64
   integer, len :: ml=100
   integer(k2) :: counter(r,c)
   character(len=ml) :: message
end type
```

```
! usage, e.g.:
type(mat rT(real32,9,9,int64,80)), target :: A
allocatable, target :: B
Class(*), pointer :: P
P => A ! P is now of dynamic type mat rT
allocate(mat crT(real64,5,5,int32,80) :: B)
P => B ! P is now of dynamic type mat crT
! unwrap polymorphism to access components
select type(P)
type is (mat_crT(real64,*,*,int32,*))
 write(*,*)'%m=',P%m
 write(*,*)'%counter=',P%counter
end select
```

- unwrap polymorphism from polymorphic object (here P) to access components
- □ argument for type-guard statement: need to specify all kind parameters (compile-time constants) and all len parameters as assumed (*-notation)



Interoperation of Fortran with C

Overview of functionality defined in





Area of semantics	within Fortran	within C	
function (procedure) call	invoke C function or interoperable Fortran procedure	invoke interoperable Fortran procedure	
main program	only one: either Fortran or C		
intrinsic data types	subset of Fortran types denoted as interoperable; not all C types are known	not all Fortran types may be known	
derived data types	special attribute enforces interoperability with C struct types	"regular" Fortran derived types not (directly) usable	
global variables	access data declared with external linkage in C	access data stored in COMMON or module variable	
dummy arguments	arrays or scalars	pointer parameters	
dummy arguments	with VALUE attribute	non-pointer parameters	

Overview continued



Dealing with I/O:

- Fortran record delimiters
- STREAM I/O already dealt with

Focus here is on: standard conforming Fortran/C interoperation

Earlier attempts

- F2C interface
- fortran.h include file
- proprietary directives

are not discussed in this course

- different concepts!
- partial semantic overlap
- procedure/function pointers

C and Fortran pointers

Semantics	within Fortran	within C
C pointer	object of type(c_ptr)	void *
C function pointer	object of type(c_funptr)	void (*)()

 module functions are provided via an intrinsic module to map data stored inside these objects to Fortran POINTERs and procedure pointers

The concept of a companion processor



Used for implementing C interoperable types, objects and functions

 it must be possible to describe function interfaces via a C prototype

Companion may be

- a C processor
- another Fortran processor supporting C interoperation
- or some other language supporting C interoperation

Note:

- different C processors may have different ABIs and/or calling conventions
- therefore not all C processors available on a platform may be suitable for interoperation with a given Fortran processor

C-Interoperable intrinsic types



Example program:

```
a module provided by
                                   the Fortran processor
program myprog
  use, intrinsic :: iso_c_binding
  integer(c_int) :: ic
  real(c float) :: rc4
  real(c double), allocatable :: a(:)
  character(c char) :: cc
                 further stuff omitted here -
                    will be shown later
  allocate(a(ic), ...)
  call my_c_subr(ic,a) \_
                                           might be implemented
                                              in Fortran or C.
                                       Will show a C implementation later
end program
```

Mapping of some commonly used intrinsic types



via KIND parameters

integer constants defined in ISO_C_BINDING intrinsic module

C type	Fortran declaration	C type	Fortran declaration	
int	integer(c_int)	char	character(len=1,kind=c_char)	
long int	integer(c_long)		may be same as kind('a')	
size_t	integer(c_size_t) may be same as c_int			
[un]signed char	integer(c_signed_char)	_Bool	logical(c_bool)	
		on x86 arch	nitecture: the same	
float	real(c_float)	type But this is not guaranteed		
double	real(c_double)			

- a negative value for a constant causes compilation failure (e.g., because no matching C type exists, or it is not supported)
- a standard-conforming processor must only support c_int
- compatible C types derived via typedef also interoperate

Calling C subprograms from Fortran: a simple interoperable interface



Assume a C prototype

```
void My_C_Subr(int, double []);
```

or double * ?

C implementation not shown

Need a Fortran interface

- explicit interface
- BIND(C,name='...') attribute
 - suppress Fortran name mangling
 - label allows mixed case name resolution and/or renaming (no label specified → lowercase Fortran name is used)
 - cannot have two entities with the same binding label

left-out bits from previous program

VALUE attribute/statement

- create copy of argument
- some limitations apply (e.g., cannot be a POINTER)

Scalar vs. array pointers

- no unique interpretation in C
- check API documentation

Functions vs. subroutines and compilation issues



C function with void result

may interoperate with a Fortran subroutine

All other C functions

may interoperate with a Fortran function

Link time considerations


- recommendation: perform linkage with Fortran compiler driver to assure Fortran RTL is linked in
- may need a special compiler link-time option if main program is in C (this is processor-dependent)

Passing arrays between Fortran and C (1)



Return to previous example:

- assume that six array elements have been allocated
- remember layout in memory: contiguous storage sequence



formal parameter double d[] uses zero-based indexing
 (C ignores any lower bound specification in the Fortran interface!)

Note:

- in a call from Fortran, a non-contiguous array (e.g. a section) may be used → will be automatically compactified (copy-in/out)
- need to do this manually in calls from C

Multidimensional arrays



Example Fortran interface

Two possible C prototypes

```
void solve_mat(double *, \
    int, int);
```

```
void solve_mat(double [*][*], \
    int, int);

C99 VLA (variable length array)
```

2nd index

Assume actual argument in call from Fortran:

```
double precision :: rm(0:1,3)
:
call solve_mat(rm, 2, 3)
```

```
1 2 3 "columap indication index
```

"column major" mapping of array indices to storage index

Accessing multidimensional array data in C



First alternative – manual mapping

example implementation:

```
void solve mat(double *d, int n1, int n2) {
  double **dmap;
  int i, k;
  dmap = (double **) malloc(n2 * sizeof(double *));
  for (i=0; i<n2; i++) {
                                                 force **dmap to contiguous
    dmap[i] = d + n1 * i;
                                                       storage layout
// now access array elements via dmap
    for (k=0; k<n1; k++) {
      dmap[i][k] = ...;
LHS is of type double
                                                              1<sup>st</sup> index
  free (dmap);
                                                                "row major"
                                                                mapping of array
                                                  1 3 5
                                                                indices to storage
                                                                index
                                            2<sup>nd</sup> index
```

Accessing multidimensional array data in C



Second alternative – C99 VLA

example implementation:

Caveat for use of ** (pointer-to-pointer):

in general this describes a non-contiguous storage sequence → cannot be used as interoperable array parameter

```
dmap[i] = (double *) malloc(...);

dmap[0]

dmap[1]

dmap[2]
```

Handling of strings (1)



- Remember: character length must be 1 for interoperability
- Example: C prototype

```
int atoi(const char *);
```

matching Fortran interface declares c_char entity an assumed size array

```
interface
  integer(c_int) function atoi(in) bind(c)
    use, intrinsic :: iso_c_binding
    character(c_char), dimension(*) :: in
  end function
end interface
```

Handling of strings (2)



Invoked by

- special exception (makes use of storage association): actual argument may be a scalar character string
- Character constants in ISO_C_BINDING with C-specific meanings

Name	Value in C
c_null_char	′\0′
c_new_line	'\n'
c_carriage_return	'\r'

most relevant subset

C Interoperation with derived types



Example:

```
use iso_c_binding
:
type, bind(c) :: dtype
  integer(c_int) :: ic
  real(c_double) :: d(10)
end type dtype
```

is interoperable with

```
typedef struct {
  int i;
  double dd[10];
} dtype_c;
```

and typed variables can be used e.g., in argument lists

Notes:

- naming of types and components is irrelevant
- bind(c) cannot have a label in this context. It cannot be specified together with sequence
- position of components must be the same
- type components must be of interoperable type

Interoperation with derived types: Restrictions



In this context, Fortran type components must not be

- pointers or allocatable
- zero-sized arrays
- type bound procedures

Fortran type must not be

extension of another type (and an interoperable type cannot itself be extended!)

C types which cannot interoperate:

- union types
- structs with bit field components
- structs with a flexible array member

Handling non-interoperable data – the question now is ...



when and how to make objects of the (non-interoperable!)
Fortran type

```
type :: fdyn
real(c_float), allocatable :: f(:)
end type fdyn
```

available within C

when and how to make objects of the analogous C type

```
typedef struct cdyn {
  int len;
  float *f;
} Cdyn;
```

available within Fortran

Case 1: Data only accessed within C



API calls are

Assumptions:

- want to call from Fortran
- but no access to type components needed within Fortran

Required Fortran interface

```
use, intrinsic :: iso c binding
interface
 type(c_ptr) function &
     cdyn create(len) bind(c,...)
   import :: c_int, c_ptr
   integer(c int), value :: len
 end function
 subroutine cdyn add(h, ...) &
                        bind(c,...)
   import :: c ptr
   type(c_ptr), value :: h
                       object of type
 end subroutine
                      c ptr requires
end interface
                    value attribute here
```

Typeless C pointers in Fortran



Opaque derived types defined in ISO_C_BINDING:

- c_ptr: interoperates with a void * C object pointer
- c_funptr: interoperates with a C function pointer.

Useful named constants:

```
c_null_ptr: C null pointer
```

```
type(c_ptr) :: p = c_null_ptr
```

c_null_funptr: C null function pointer

Logical module function that checks pointer association:

- c_associated(c1[,c2])
- value is .false. if c1 is a C null pointer or if c2 is present and points to a different address. Otherwise, .true. is returned
- typical usage:

```
type(c_ptr) :: res

res = get_my_ptr( ... )
if (c_associated(res)) then
    :! do work with res
else
    stop 'NULL pointer produced by get_my_ptr'
end if
```

Case 1 (cont'd): Client usage



```
use, intrinsic :: iso_c_binding
:
type(c_ptr) :: handle
:
handle = cdyn_create(5_c_int)
if (c_associated(handle)) then
    call cdyn_add(handle,...)
end if
call cdyn_destroy(handle)
all memory
management done in C
```

Typeless "handle" object

 because objects of (nearly) any type can be referenced via a void *, no matching type declaration is needed in Fortran

Design problem:


 no disambiguation between different C types is possible → loss of type safety

possible solution

Setting up a mapping between a Fortran object and a C pointer



Module ISO_C_BINDING provides module procedures



 pointer association (blue arrow) is set up as a result of their invocation (green arrows)

Two scenarios are covered by c_loc/c_f_pointer



- 1. Fortran object is of interoperable type and type parameters:
 - variable with target attribute,
 - or allocated variable with target attribute, non-zero length,
 - or associated scalar pointer

in scenario 1, the object might also have been created within C (Fortran target then is anonymous).

In any case, the data can be accessed from C.

- 2. Fortran object is a noninteroperable, nonpolymorphic scalar without length type parameters:
 - non-allocatable, non-pointer variable with target attribute,
 - or an allocated allocatable variable with target attribute,
 - or an associated pointer.

nothing can be done with such an object within C

Case 2: Data accessed only within Fortran



Fortran Library definition

- noninteroperable derived data type
- provide an interoperable constructor written in Fortran

```
type (c_ptr) function &
    fdyn_create(len) bind(c,...)
    integer(c_int), value :: len
    type(fdyn), pointer :: p
    allocate(p)
    allocate(p%f(len))
    fdyn_create = c_loc(p)
    end function
end module mylib
```

Pointer goes out of scope

- but target remains reachable via function result
- C prototype:

```
void *fnew_stuff(int);
```

Case 2 (cont'd): Retrieving the data



Client code in C:

```
void *fhandle;
int len = 5;
do not try to
dereference handle
except for NULL check

fhandle = Fdyn_create(len);

Fdyn_print(fhandle);
```

- can have multiple handles to different objects at the same time (thread-safeness)
- again no matching type needed on client
- require Fortran implementation of Fdyn print()

... here it is:

... and must not forget to

- implement "destructor" (in Fortran)
- and call it (from C or Fortran) for each created object

to prevent memory leak

Warning on inappropriate use of c_loc() and c_f_pointer()



With these functions,

- it is possible to subvert the type system (don't do this!)
 (push in object of one type, and extract an object of different type)
- it is possible to subvert rank consistency (don't do this!)
 (push in array of some rank, and generate a pointer of different rank)

Implications:

- implementation-dependent behaviour
- security risks in executable code

Recommendations:

- use with care (testing!)
- encapsulate use to well-localized code
- don't expose use to clients if avoidable

Case 3: Accessing C-allocated data in Fortran



- We haven't gone the whole way towards fully solving the problem
 - won't actually do so in this talk
- Return to Case 1:

```
typedef struct Cdyn {
  int len;
  float *f;
} Cdyn;
```

```
void Cdyn_print(Cdyn *);
```

- and implement the function with above C prototype in Fortran
- → need read and/or write access to data allocated within the Cdefined structure
- allocation is performed as described in Case 1

Case 3 (cont'd): Fortran implementation



Required type definition:

Notes:

- note the intent(in) for this (refers to association of c_ptr; the referenced data can be modified)
- scenario 1 applies for c_f_pointer usage

Implementation:

```
subroutine cdyn_print(this) bind(c,name='Cdyn_print')
  type(cdyn), intent(in) :: this
  real(c_float), pointer :: cf(:)
! associate array pointer cf with this%f
  call c_f_pointer( this%f, cf, [this%len] )
! now do work with data pointed at by this%f
  write(*,fmt=...) cf
end subroutine cdyn_print
```

following now: Exercise session 10

Procedure arguments and pointers (1)



Procedure argument: a function pointer in C

```
    could have a fixed or variable interface
```

Matched by interoperable Fortran interface

```
real(c_double) function integrate(a, b, par, fptr) bind(c)
  real(c_double), value :: a, b
  type(c_ptr), value :: par
  type(c_funptr), value :: fptr
end function
a C function pointer
```

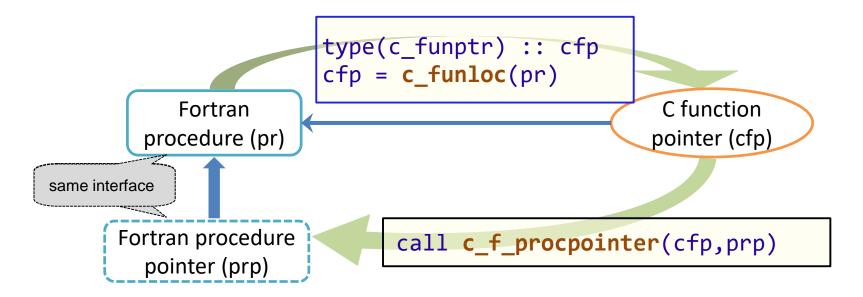
■ Note:

 an interface with a Fortran procedure dummy argument is not interoperable (even if the dummy procedure has the BIND(C) attribute)

Setting up a mapping between a Fortran procedure and a C procedure pointer



Module ISO_C_BINDING provides module procedures



- input for c_funloc must be an interoperable Fortran procedure;
 can also be an associated procedure pointer
- pointer association (blue arrow) is set up as a result of their invocation (green arrows)

Procedure arguments and pointers (2)



Assuming the function interface

note the consistent interface

```
real(c_double) function f(x, par) bind(c)
  real(c_double), value :: x
  type(c_ptr), value :: par
end function
```

the invocation reads

```
type(c_funptr) :: fp
fp = c_funloc(f)
res = integrate( a, b, par, fp )
```

or, more concisely

```
res = integrate( a, b, par, c_funloc(f) )
```

Another procedure pointer example



C function pointer used as type component

```
typedef struct {
  double (*f)(double, void *);
  void *par;
} ParFun;
```

Matching type definition in Fortran:

requires use of component of type c_funptr

```
type, BIND(C) :: parfun
    type(c_funptr) :: f
    type(c_ptr) :: par
end type
```

Invoking the C-associated procedure from Fortran



Example:

```
type(parfun) :: o_pf
o_pf%f = c_funloc(my_function)
o_pf%par = c_loc(...)
```

```
ParFun *o_pf;
o_pf->f = my_function;
o_pf->par = (void *) ...;
```

 where my_function should have the same interface in Fortran and C, respectively.

```
type(parfun) :: o_pf
type(c_ptr) :: par
procedure(my_function), pointer :: pf

: ! initialize o_pf, par within C or Fortran
call c_f_procpointer(o_pf%f, pf)
y = pf(2.0_dk, par)
```

Interoperation of global data (1): COMMON blocks



Defining C code:

```
struct coord{
  float xx, yy
};
struct coord csh;
```

- do not place in include file
- reference with external in other C source files

Mapping Fortran code

```
real(c_float) :: x, y
common /csh/ x, y
bind(c) :: /csh/
```

- BIND statement (possibly with a label) resolves to the same linker symbol as defined in C → same memory address
- memory layout may be different as for "traditional" sequence association

Interoperation of global data (2): Module variables



Defining C code:

```
int ic;
float Rpar[4];
```

- do not place in include file
- reference with external in other C source files

Mapping Fortran code:

```
module mod_globals
  use, intrinsic :: iso_c_binding

integer(c_int), bind(c) :: ic
  real(c_float) :: rpar(4)
  bind(c, name='Rpar') :: rpar
end module
```

 either attribute or statement form may be used

Global binding can be applied to objects of interoperable type and type parameters.

Variables with the ALLOCATABLE/POINTER attribute are not permitted in this context.

Enumeration



Set of integer constants

only for interoperation with C

```
enum, bind(c)
  enumerator :: red=4, blue=9
  enumerator :: yellow
end enum
```

- integer of same kind as used in C enum
- value of yellow is 10
- not hugely useful

Extension of interoperability with C in





- Preliminary specification (2012): "Mini-standard" ISO/IEC TS 29113
- Motivations:
 - enable a standard-conforming MPI (3.1) Fortran interface
 - permit C programmers (limited) access to "complex" Fortran objects

Area of semantics	within Fortran	within C
dummy argument POINTER or ALLOCATABLE	assumed shape/length or deferred shape/length	pointer to a descriptor
dummy argument	assumed rank	pointer to a descriptor
dummy argument	assumed type	either void * or pointer to a descriptor
dummy argument	OPTIONAL attribute no VALUE attribute permitted	use a NULL actual or check formal for being NULL
dummy argument of type c_ptr or c_funptr	non-interoperable data or procedure	handle only, no access to data or procedure
non-blocking procedures	ASYNCHRONOUS attribute	not applicable

Accessing Fortran infrastructure from C: the source file ISO_Fortran_binding.h



Example Fortran interface

e Matching C prototype

```
subroutine process_array(a) BIND(C)
  real(c_float) :: a(:,:)
  assumed
  shape
```

- Implementation of procedure might be in C or in Fortran
- For an implementation in C, the header provides access to
 - type definition of descriptor (details upcoming ...)
 - macros for type codes, error states etc.
 - prototypes of library functions that generate or manipulate descriptors
- Reserved namespace: CFI_
- Within a single C source file,
 - binding is only possible to one given Fortran processor (no binary compatibility!)

Members of the C descriptor




Exposes internal structure of Fortran objects

not meant for tampering



→ please use API



Type and attribute members



Type code macros

most commonly used:

Name	C type
CFI_type_int	int
CFI_type_long	long int
CFI_type_size_t	size_t
CFI_type_float	float
CFI_type_double	double
CFI_type_Bool	_Bool
CFI_type_char	char
CFI_type_cptr	void *
CFI_type_struct	Interoperable C structure
CFI_type_other (<0)	Not otherwise specified

Attribute of dummy object

Name	
CFI_attribute_allocatable	
CFI_attribute_pointer	
CFI_attribute_other	
e.g., assumed shape or length	

 Beware: attribute value of actual must match up exactly with that of dummy (different from Fortran)
 → may need to create descriptor copies

typically, non-interoperable data

Using the descriptor to process array elements (1)



Fortran reference loop within process_array():

Remember: "a" represents a rank-2 array of assumed shape

C implementation variant 1:

ordering of dimensions as in Fortran

```
for (k = 0; k < a->dim[1].extent; k++) {
   for (i = 0; i < a->dim[0].extent; i++) {
      CFI_index_t subscripts[2] = { i, k };
      ... = *((float *) CFI_address( a, subscripts )) * ...;
   }
}
```

- CFI_address() returns (void *) address of array element indexed by specified (valid!) subscripts
- dim[].lower_bound will be needed for pointer/allocatable objects

Using the descriptor to process array elements (2)



C implementation variant 2:

start out from beginning of array

```
char *a_ptr = (char *) a->base_addr;
```

and use pointer arithmetic to process it:

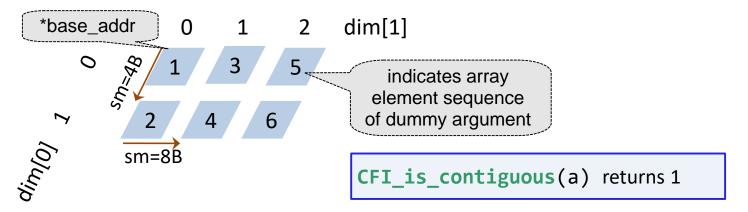
```
char *a_aux;
for (k = 0; k < a->dim[1].extent; k++) {
    a_aux = a_ptr;
    for (i = 0; i < a->dim[0].extent; i++) {
        ... = *((float *) a_ptr) * ...;
        a_ptr += a->dim[0].sm;
        pointer to Fortran
        element a(i,k)
    a_ptr = a_aux + a->dim[1].sm;
}
```

- non-contiguous arrays require use of stride multipliers (next slide illustrates why)
- stride multipliers in general may not be an integer multiple of the element size → always process in units of bytes


Memory layouts for assumed shape objects



Actual argument is a complete array (0:1,3)



Actual argument is an array section (0::2,1::3) of (0:2,9)



Creating a Fortran object within C



- May be necessary to invoke a Fortran procedure from C
- Step 1: create a descriptor

```
cFI_CDESC_T(2) A; use macro to establish needed storage; maximum rank must be specified as parameter

CFI_cdesc_t *a = (CFI_cdesc_t *) &A; cast to properly typed pointer
```

Step 2: establish object's properties

prototype of function to be used for this is

- many usage patterns
- if fully defined, result is always a contiguous object
- function result is an error indicator (CFI_SUCCESS → OK)

Table of return value macros



Macro name	Explanation of error
CFI_SUCCESS	No error detected.
CFI_ERROR_BASE_ADDR_NULL	The base address member of a C descriptor is a null
	pointer in a context that requires a non-null pointer
	value.
CFI_ERROR_BASE_ADDR_NOT_NULL	The base address member of a C descriptor is not a null
	pointer in a context that requires a null pointer value.
CFI_INVALID_ELEM_LEN	The value supplied for the element length member of a C
	descriptor is not valid.
CFI_INVALID_RANK	The value supplied for the rank member of a C descriptor
	is not valid.
CFI_INVALID_TYPE	The value supplied for the type member of a C descriptor
	is not valid.
CFI_INVALID_ATTRIBUTE	The value supplied for the attribute member of a C
	descriptor is not valid.
CFI_INVALID_EXTENT	The value supplied for the extent member of a CFI_dim_t
	structure is not valid.
CFI_INVALID_DESCRIPTOR	A general error condition for C descriptors.
CFI_ERROR_MEM_ALLOCATION	Memory allocation failed.
CFI_ERROR_OUT_OF_BOUNDS	A reference is out of bounds.

Example: create rank 2 assumed shape array



```
#define DIM1 56
#define DIM2 123
CFI CDESC T(2) A; /* 2 is the minimum value needed */
CFI cdesc t *a = (CFI cdesc t *) &A;
CFI index t extents[2] = { DIM1, DIM2 };
                         /* shape of rank 2 array */
float *a ptr = (float *) malloc(DIM1*DIM2*sizeof(float));
                         /* heap allocation within C */
                         /* initialize values of *a ptr */
CFI_establish( a, (void *) a ptr,
               CFI attribute other,
               CFI type float,
               0, /* elem_len is ignored here */
                            /* rank as declared in Fortran */
               2,
               extents );
                         /* have a fully defined object now */
process array(a);
                        /* object becomes invalid */
free(a_ptr);
```

Allocatable objects



Typically only needed if Fortran API defines a "factory":

```
type, bind(c) :: qbody
                                       typedef struct {
 real(c float) :: mass
                                         float mass;
 real(c float) :: position(3)
                                         float position[3];
end type
                                         qbody;
                                                        matching C
interface
                                                      struct definition
  subroutine qbody_factory(this, fname) bind(c)
    type(qbody), allocatable, intent(out) :: this(:,:)
    character(c char, len=*), intent(in) :: fname
  end subroutine
end interface
```

matching C prototype:

descriptor corresponds to assumed length character

```
void qbody_factory(CFI_cdesc_t *, CFI_cdesc_t *)
```

Example: create an allocatable entity and populate it



```
char fname ptr[] = "InFrontOfMyHouse.dat";
CFI cdesc t *pavement =
          (CFI cdesc t *) malloc(sizeof(CFI CDESC T(2)));
CFI_cdesc t *fname =
          (CFI cdesc t *) malloc(sizeof(CFI CDESC T(0)));
                                must start out unallocated
CFI establish( pavement, NULL, CFI_attribute_allocatable,
                CFI type struct.
                sizeof(qbody), /* derived type object size */
                2, NULL ); shape is deferred
CFI establish( fname, fname ptr, CFI attribute other,
                CFI type char,
                strlen(fname_ptr), /* a char has one byte */
                0, NULL );
qbody factory ( pavement, fname ); /* object is created */
                                    /* process pavement */
CFI_deallocate( pavement ); _____
free(pavement); free(fname);
                                     no auto-deallocation of objects allocated in C
```

An implementation of qbody_factory() in C



Feasible because of supplied function CFI_allocate():

 last argument is an element length, which is ignored unless the type member is CFI_type_char. In the latter case, its value becomes the element length of the allocated deferred-length (!) string.

following now: Exercise session 11

Handling Fortran POINTERs within C



Anonymous target

- create descriptor with CFI_attribute_pointer, then apply CFI_allocate()/CFI_deallocate()
- Point at an existing target

```
real(c_float), TARGET :: t(:)
real(c_float), POINTER :: p(:)
p(3:) => t
```

- t must describe a fully valid object
- p must be an established descriptor with
 CFI_attribute_pointer and for the same type as t.



Beware: No compile-time type safety is provided.

Certain inconsistencies may be diagnosed at run time

→ check return value of CFI_setpointer()

Creating subobjects in C (1)



Assumption:

- arr describes an assumed-shape rank 3 array
- Create a descriptor for the section arr(3:,4,::2)

```
CFI cdesc t *section =
          (CFI cdesc t *) malloc(sizeof(CFI CDESC T(2)));
CFI index t lower_bounds[3] = \{ 2, 3, 0 \};
CFI index t upper bounds[3] =
          { arr->dim[0].extent - 1, 3, arr->dim[2].extent - 1 };
CFI index t strides[3] = \{ 1, 0, 2 \};
                            zero stride indicates a subscript.
                 For this dimension, lower and upper bounds must be equal.
CFI establish( section, NULL, CFI_attribute_other,
                arr->type, arr->elem len, 2, NULL );
                /* section here is an undefined object */
CFI_section( section, arr, lower_bounds, upper_bounds, strides );
                /* now, section is defined */
```

Creating subobjects in C (2)



Type component selection

- pavement(:)%position(1) from the type(qbody) object pavement
- a rank-2 array of intrinsic type real(c_float)

displacement

sm=16B

pos_1

pavement(1:4,1:1)

Assumed rank dummy argument



Enables invocation of appropriately declared object

```
subroutine process_allranks(ar, ...)
  real :: ar(..)
  ...
  write(*,*) RANK(ar)
end subroutine
```

ar cannot (currently) be referenced or defined within Fortran.

However, some intrinsics can be invoked.

with arrays of any rank, or even a scalar:

```
real :: xs, x1(4), x2(ndim, 4)

call process_allranks(xs, ...) scalar rank 1 rank 1 rank 2
```

avoid need for writing many specifics for a named interface



Assuming the procedure interface is made BIND(C):

descriptor always contains well-defined rank information

```
void process_allranks(CFI_cdesc_t *ar, ...) {
   switch( ar->rank )
   case 1:
        ... /* process single loop nest */
   case 2:
        ... /* process two nested loops */
   default:
      printf("unsupported rank value\n");
      exit(1);
   }
}
```

 deep loop nests can be avoided for contiguous objects, but the latter is not assured

Assumed size actual argument



Special case:

size of (contiguous) assumed-size object is not known

```
real :: x2(ndim, *)

call process_allranks(x2, ..., ntot)

should specify size separately
```

A descriptor with following properties is constructed:

- SIZE(ar,DIM=RANK(ar)) has the value -1
- UBOUND(ar,DIM=RANK(ar)) has the value UBOUND(ar,DIM=RANK(ar)) - 2

Assumed type dummy arguments



Declaration with TYPE(*)

- an unlimited polymorphic object → actual argument may be of any type
- dynamic type cannot change → no POINTER or ALLOCATABLE attribute is permitted

Corresponding object in interoperating C call:

two variants are possible

TYPE(*), DIMENSION(*) :: obj

TYPE(*) :: obj

TYPE(*), DIMENSION(:) :: obj

TYPE(*), DIMENSION(..) :: obj

CFI_cdesc_t *obj

Example: direct interoperation with MPI_Send



C prototype as specified in the MPI standard

```
int MPI_Send( const void *buf, int count, MPI_Datatype datatype,
    int dest, int tag, MPI_Comm comm );
```

Matching Fortran interface:

- assumes interoperable types MPI_Datatype etc.
- actual argument may be array or scalar
- non-contiguous actuals are compactified

array temps are a problem for non-blocking calls

MPI-3 Fortran interface for MPI_send



```
SUBROUTINE MPI_Send( buf, count, datatype, dest, & tag, comm, ierror )

TYPE(*), DIMENSION(..), INTENT(IN) :: buf
INTEGER, INTENT(IN) :: count, dest, tag
TYPE(MPI_Datatype), INTENT(IN):: datatype
TYPE(MPI_Comm), INTENT(IN) :: comm
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END FUNCTION MPI_Send
```

Invocation of MPI_Send

 now possible also with array section actual arguments without need for copy-in/out

Could add BIND(C) to the interface for a C implementation

- assuming int matches default Fortran integer
- the MPI standard doesn't do this, though

C implementation of MPI_Send() wrapper



```
void mpi send( CFI cdesc t *buf, int count, MPI Datatype datatype,
               int dest, int tag, MPI_Comm comm, int *ierror ) {
   int ierror local;
   MPI Datatype disc type;
   if ( CFI is contiguous( buf ) ) {
     ierror_local = MPI_Send( buf->base_addr, count, datatype,
                    dest, tag, comm );
   } else {
     ... /* use descriptor information to construct disc type
             from datatype (e.g. via MPI_Type_create_subarray) */
     ierror_local = MPI_Send( buf->base_addr, count, disc_type,
                    dest, tag, comm );
     ... /* clean up disc type */
   if (ierror != NULL) *ierror = ierror local;
```

Automatized translation of C include files to Fortran interface modules



- Requires a specialized tool
 - for example, Garnet Lius LLVM-based tool, see https://github.com/Kaiveria/h2m-Autofortran-Tool
- C include files can have stuff inside that is not covered by interoperability
 - only a subset can be translated
- Topic goes beyond the scope of this course

Final remarks



Interoperation with C++

- no direct interoperation with C++-specific features is possible
- you need to write C-like bridge code
- declare C-style functions "extern C" in your C++ sources
- explicit linkage of C++ libraries will be needed if the Fortran compiler driver is used for linking

Vararg interfaces

- are not interoperable with any Fortran interface
- you need to write glue code in C



Shared Libraries and Plug-ins

What is a shared library?



Executable code in library

- is shared between all programs linked against the library (instead of residing in the executable)
- this does not apply to data entities

Advantages:

- save memory space
- save on access latency
- bug fixes in library code do not require relinking the application

Disadvantages:

- higher complexity in handling the build and packaging of applications
- (need to distribute shared libraries together with the linked application)
- not supported (in analogous manner) on all operating environments
- (will focus on ELF-based Linux in this talk)
- special compilation procedure is required for library code

ELF → executable and linkable format

see http://en.wikipedia.org/wiki/Executable_and_Linkable_Format for details

Compatibility issues



Causes of incompatibility

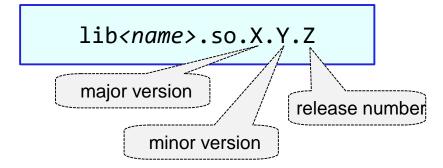
- change function interface
- remove function interface
- (adding a new function is no problem)
- change in type definitions or data items
- (exception: storage association is preserved)
- changes in function behaviour

If one of these happens,

 a mechanism is available to indicate the library is not compatible

Concept of soname

 naming scheme for shared libraries



soname will typically be

 and the latest version of the library with the same soname should be picked up at linkage

Example (1): building a library



Assume you have

a source file mylib.f90

separate steps

Implementation-dependent

options

Compiler	Compilation	Linkage
Intel ifort	-fPIC	-shared -W1,-soname lib <name>.so.X</name>
Gfortran	-fPIC	-shared -W1,-soname=lib <name>.so.X</name>
PGI pgf90	-fPIC	-shared -W1,-soname=lib <name>.so.X</name>
NAG nagfor	-PIC	-shared -W1,-soname=lib <name>.so.X</name>
IBM xIf	-G -qpic=big	-qmkshrobj
Cray ftn	?	?

Example (Intel compiler):

```
ifort -c -fPIC mylib.f90
ifort -o libmylib.so.1.0.0 -shared -Wl,-soname libmylib.so.1 mylib.o
```

Add symbolic links (Linux)

```
ln -s libmylib.so.1.0.0 libmylib.so.1
ln -s libmylib.so.1 libmylib.so
```

Linking against the library and running the executable



Linux linkage:

- specify directory where shared library resides
- specify shorthand for library name

Execute binary

- set library path
- execute as usual

```
export LD_LIBRARY_PATH=\
$HOME/lib:$LD_LIBRARY_PATH
./myprog.exe
libmylib.so lives there
```

ifort -o myprog.exe myprog.f90 -L../lib -lmylib

- note: if both a static and a shared library are found, the shared library will be used by default
- there usually exist compiler switches which enforce static linking

- note: /etc/ld.so.conf
 contains library paths which are always searched
- there usually exist possibilities to hard-code the library path into the executable

don't need to set LD_LIBRARY_PATH
in these two cases

Special linkage options



- The -W1, option can be used to pass options to the linker
- Example 1:
 - want to specify that a certain library -lspecial should be linked statically, others dynamically
 - this is not uniquely resolvable from the library specification if both static and dynamic versions exist!

Example 2: hard-code path into binary

avoids the need to set LD_LIBRARY_PATH before execution

Dynamic loading (1)



Supported by C library:

- open a shared library at run time
- extract a symbol (function pointer)
- execute function
- close shared library

man 3p dlopen / dlsym / dlclose

From Fortran

- usable via C interoperability and pointers to procedures
- implement plug-ins

Small Fortran module dlfcn

- type definition dlfcn_handle
- procedures
 dlfcn_open(),
 dlfcn_symbol(),
 dlfcn_close()
- Note: the result of dlfcn_symbol() is of type c_funptr to enable conversion to an explicit interface procedure pointer
- constants required for dlfcn_open() mode

Dynamic loading (2): An example program



```
use dlfcn
implicit none
abstract interface
  subroutine set(i) bind(c)
    integer, intent(inout) :: i
  end subroutine set
end interface
integer :: i, istat
type(dlfcn handle) :: h
type(c_funptr) :: cp
procedure(set), pointer :: fp
h = dlfcn open('./libset1.so', &
                RTLD NOW)
cp = dlfcn symbol(h, 'set val')
call c f procpointer(cp, fp)
i = 1
                                procedure
call fp(i)
                                  name
istat = dlfcn close(h)
```

Shared library libset1.so:

BIND(C) procedure

Module procedure:

explicit name mangling needed

```
h = dlfcn open('./libset2.so', &
                RTLD NOW)
! at most one line valid
cp = dlfcn_symbol(h, &
                                 gfortran
     ' s MOD set val')
cp = dlfcn symbol(h, &
                                   ifort
     's mp set val ')
cp = dlfcn_symbol(h, &
                                    xlf
     ' s NMOD set val')
cp = dlfcn symbol(h, &
                                NAG, g95
     's MP set val')
call c_f_procpointer(cp, fp)
i = 1
                                   OK with
call fp(i)
                                   TS 29113
istat = dlfcn close(h)
```

nm libset2.so | grep -i set_val

Wrapping up



This concludes the LRZ part of this course

Following now: Exercise session 12